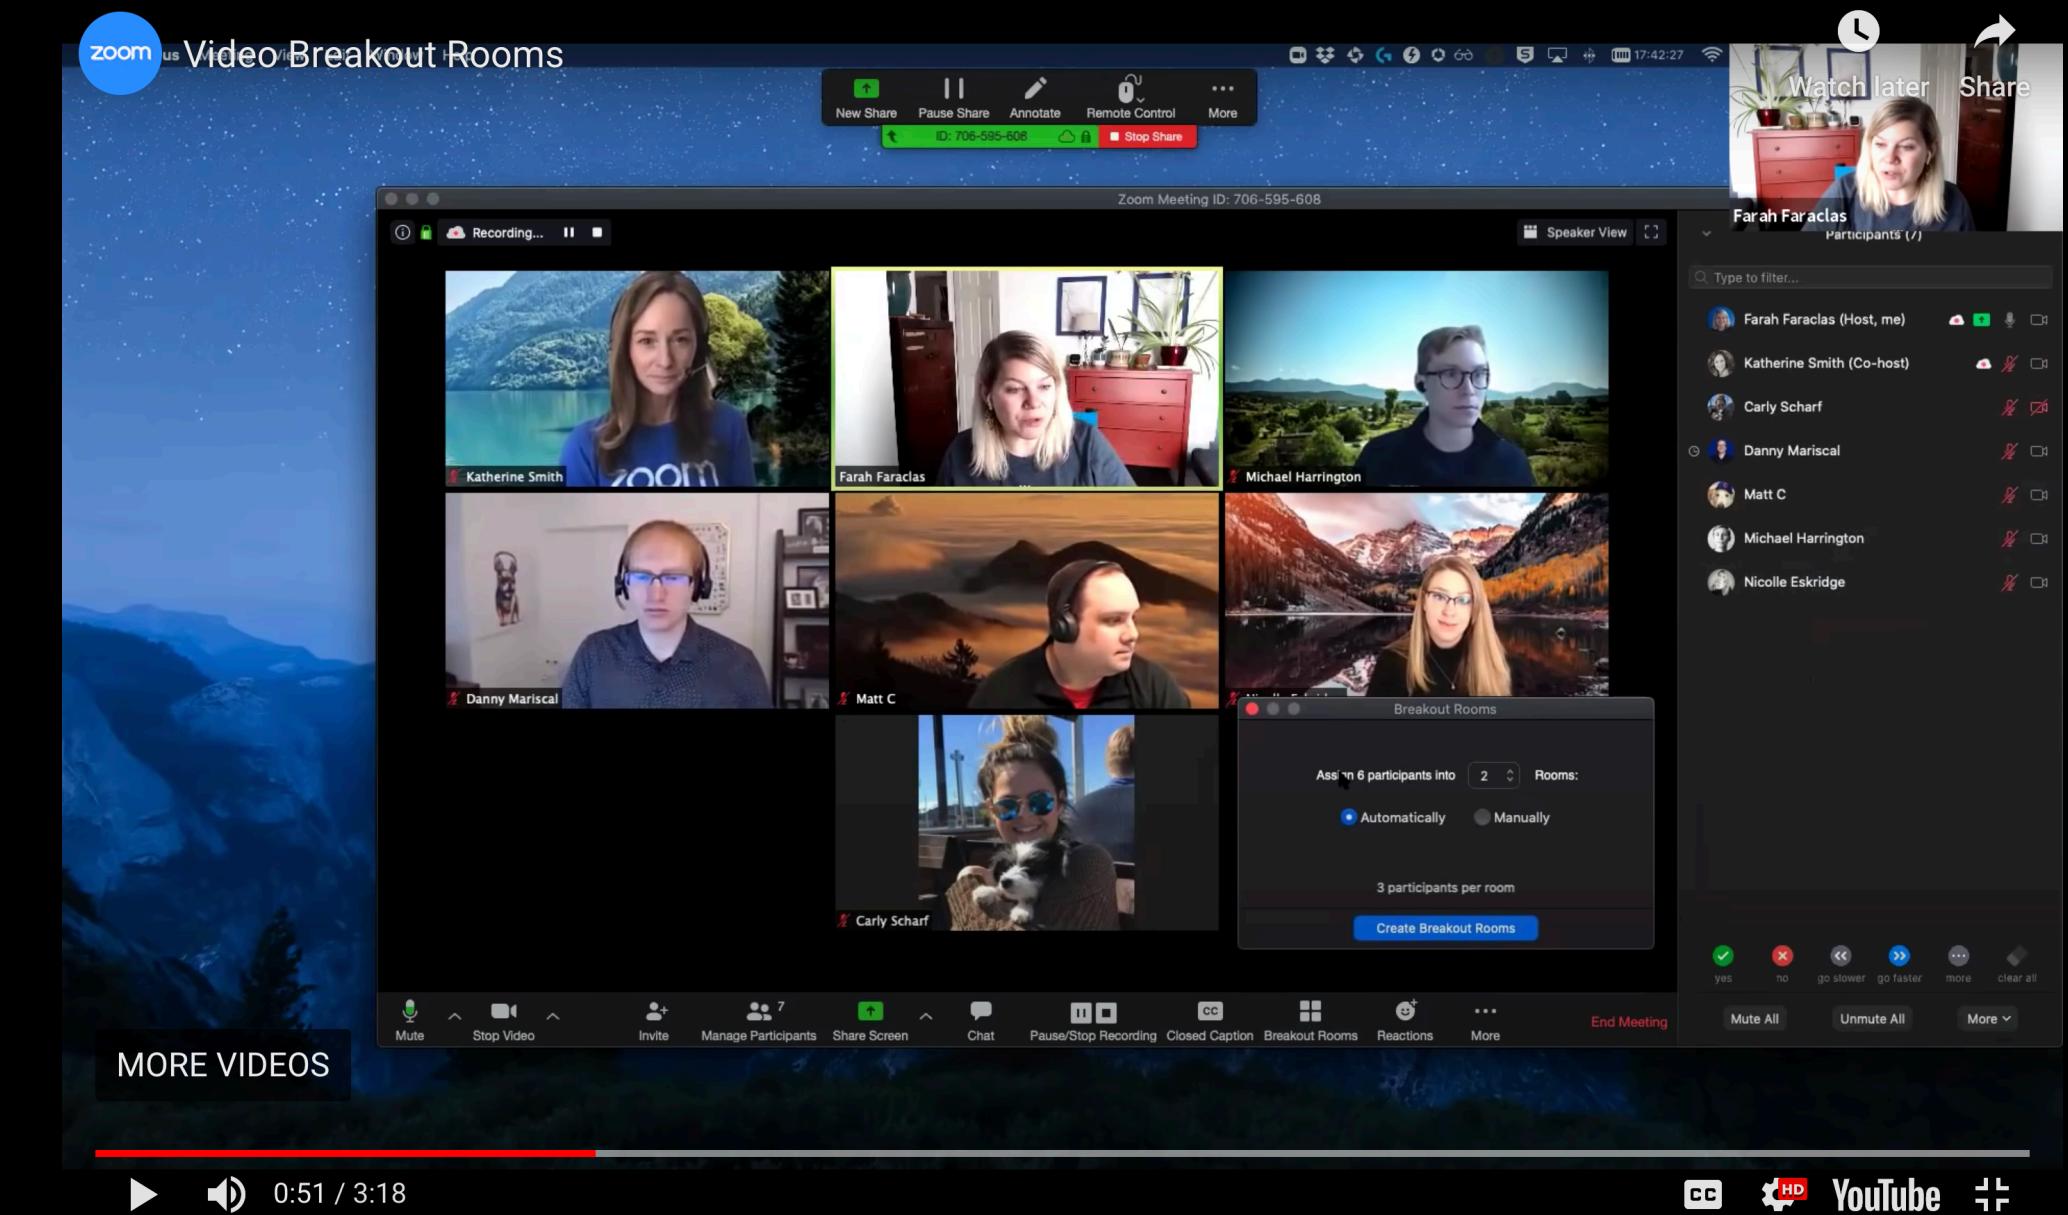


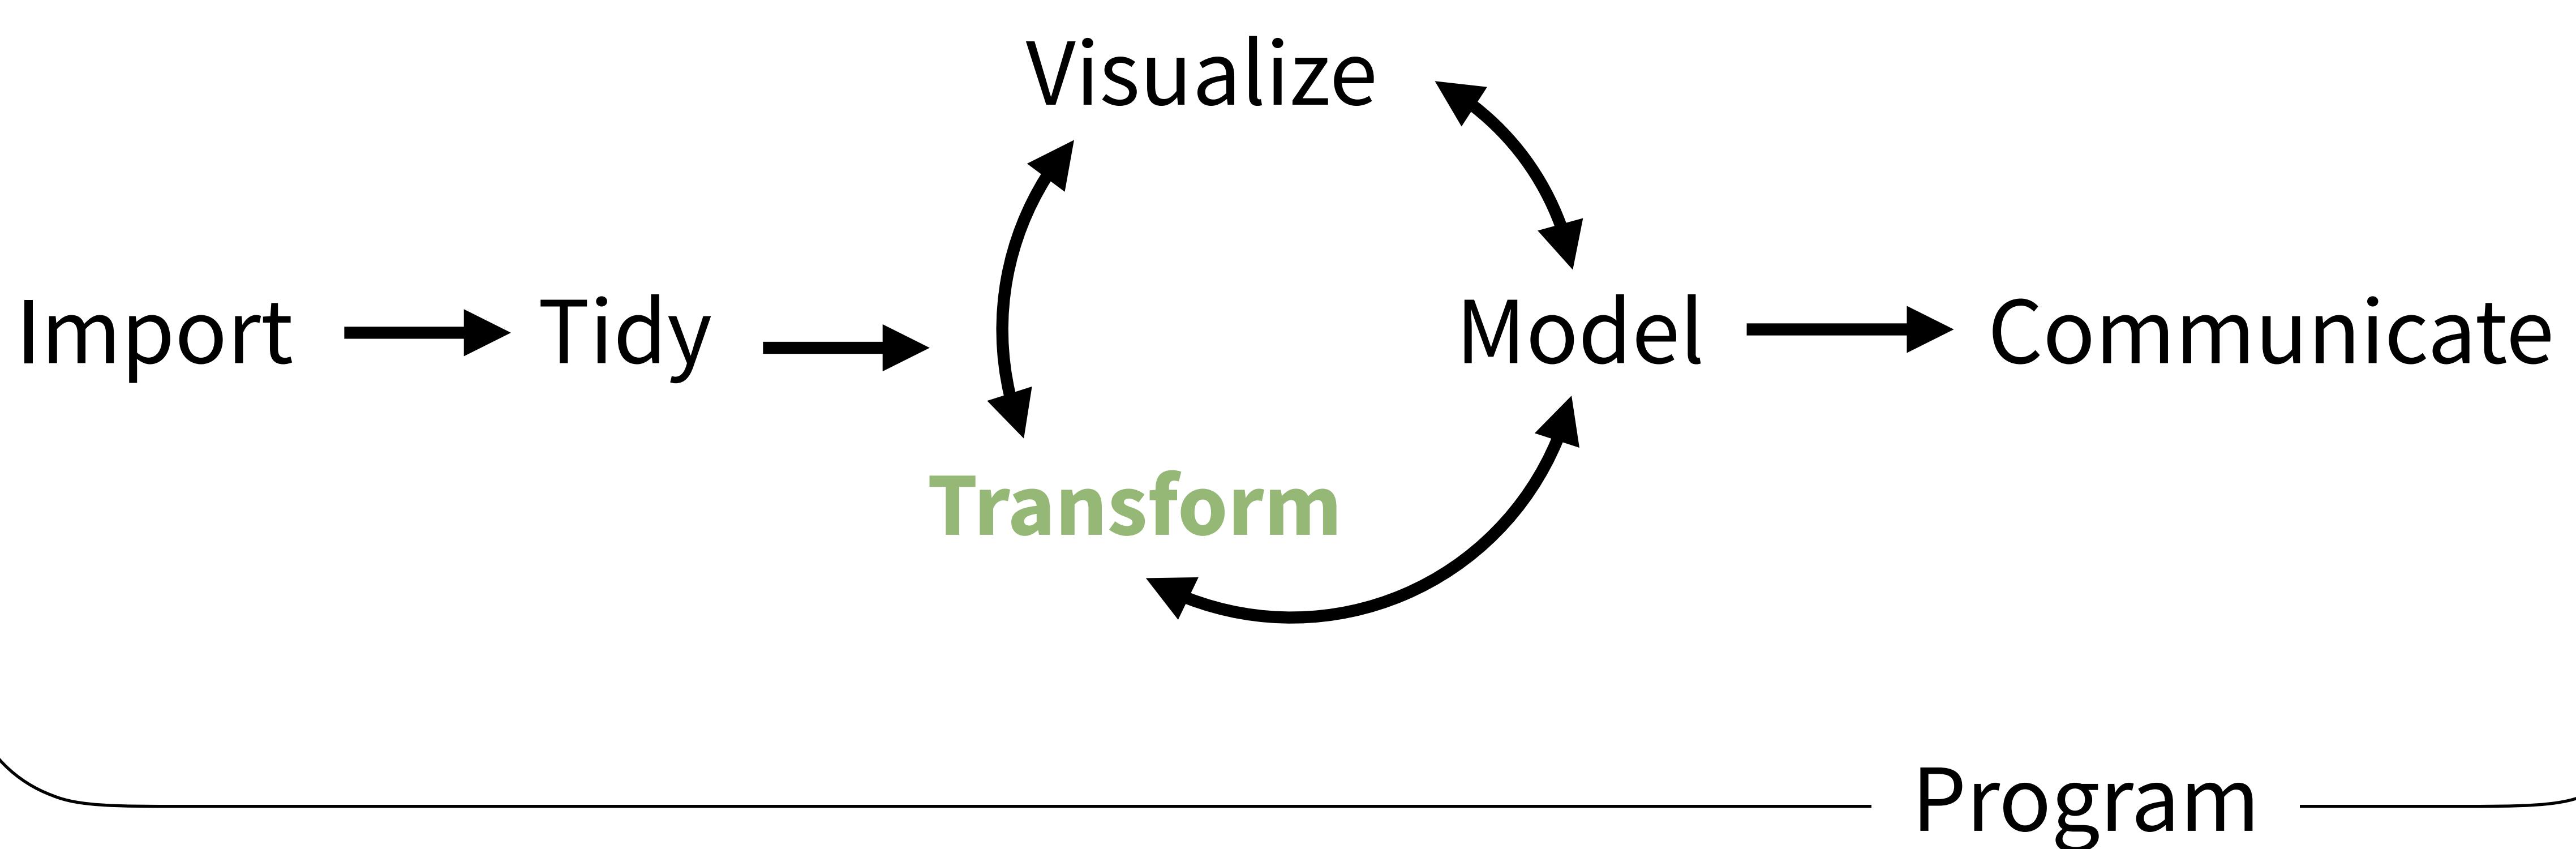
Data Types in the Tidyverse

Exercise: What do you already know?

- In the Chat Window type:
 1. What types of data do you need to work with / manipulate?
 2. What functions / packages do you use to work with them?



(Applied) Data Science



Data Types

Main Ideas

Notes

Notes Form

Please Take out

Data types with



Open Data-Types.Rmd

Quiz

What types of data are in this data set?

	time_hour	name	air_time	distance	day	delayed
1	2013-01-01 05:00:00	United Air Lines Inc.	13620s (~3.78 hours)	1400	Tuesday	TRUE
2	2013-01-01 05:00:00	United Air Lines Inc.	13620s (~3.78 hours)	1416	Tuesday	TRUE
3	2013-01-01 05:00:00	American Airlines Inc.	9600s (~2.67 hours)	1089	Tuesday	TRUE
4	2013-01-01 05:00:00	JetBlue Airways	10980s (~3.05 hours)	1576	Tuesday	FALSE
5	2013-01-01 06:00:00	Delta Air Lines Inc.	6960s (~1.93 hours)	762	Tuesday	FALSE
6	2013-01-01 05:00:00	United Air Lines Inc.	9000s (~2.5 hours)	719	Tuesday	TRUE
7	2013-01-01 06:00:00	JetBlue Airways	9480s (~2.63 hours)	1065	Tuesday	TRUE
8	2013-01-01 06:00:00	ExpressJet Airlines Inc.	3180s (~53 minutes)	229	Tuesday	FALSE
9	2013-01-01 06:00:00	JetBlue Airways	8400s (~2.33 hours)	944	Tuesday	FALSE
10	2013-01-01 06:00:00	American Airlines Inc.	8280s (~2.3 hours)	733	Tuesday	TRUE
11	2013-01-01 06:00:00	JetBlue Airways	8940s (~2.48 hours)	1028	Tuesday	FALSE

Logicals



Logicals

R's data type for boolean values (i.e. TRUE and FALSE).

```
typeof(TRUE)  
## "logical"
```

```
typeof(FALSE)  
## "logical"
```

```
typeof(c(TRUE, TRUE, FALSE))  
## "logical"
```



```
flights %>%  
  mutate(delayed = arr_delay > 0) %>%  
  select(arr_delay, delayed)
```

arr_delay <dbl>	delayed <lgl>
11	TRUE
20	TRUE
33	TRUE
-18	FALSE
-25	FALSE
12	TRUE
19	TRUE
-14	FALSE
-8	FALSE
8	TRUE



```
flights %>%  
  mutate(delayed = arr_delay > 0) %>%  
  select(arr_delay, delayed)
```

arr_delay <dbl>	delayed <lgl>
11	TRUE
20	TRUE
33	TRUE
-18	FALSE
-25	FALSE
12	TRUE
19	TRUE
-14	FALSE
-8	FALSE
8	TRUE

Can we compute
the proportion of
NYC flights that
arrived late?



Most useful skills

1. Math with logicals



Math

When you do math with logicals, **TRUE becomes 1** and
FALSE becomes 0.



Math

When you do math with logicals, **TRUE becomes 1** and **FALSE becomes 0**.

- The **sum** of a logical vector is the **count of TRUEs**

```
sum(c(TRUE, FALSE, TRUE, TRUE))
```

```
## 3
```



Math

When you do math with logicals, **TRUE becomes 1** and **FALSE becomes 0**.

- The **sum** of a logical vector is the **count of TRUEs**

```
sum(c(TRUE, FALSE, TRUE, TRUE))  
## 3
```

- The **mean** of a logical vector is the **proportion of TRUEs**

```
mean(c(1, 2, 3, 4) < 4)  
## 0.75
```



Your Turn 1

Use flights to create **delayed**, a variable that displays whether a flight was delayed (`arr_delay > 0`).

Then, remove all rows that contain an NA in **delayed**.

Finally, create a summary table that shows:

1. How many flights were delayed
2. What proportion of flights were delayed



```
flights %>%  
  mutate(delayed = arr_delay > 0) %>%  
  drop_na(delayed) %>%  
  summarise(total = sum(delayed), prop = mean(delayed))  
## # A tibble: 1 × 2  
##   total      prop  
##   <int>      <dbl>  
## 1 133004 0.4063101
```



Strings

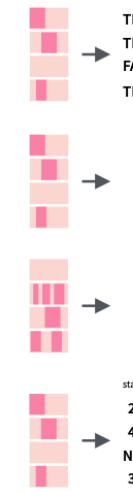
A large, semi-transparent watermark of the R logo is positioned in the bottom right corner. The logo consists of a circular emblem with the letters "R" inside.

String manipulation with stringr :: CHEAT SHEET

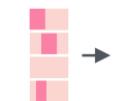
The `stringr` package provides a set of internally consistent tools for working with character strings, i.e. sequences of characters surrounded by quotation marks.



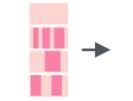
Detect Matches



`str_detect(string, pattern)` Detect the presence of a pattern match in a string.
`str_detect(fruit, "a")`



`str_which(string, pattern)` Find the indexes of strings that contain a pattern match.
`str_which(fruit, "a")`

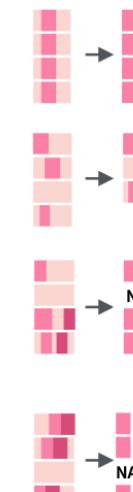


`str_count(string, pattern)` Count the number of matches in a string.
`str_count(fruit, "a")`



`str_locate(string, pattern)` Locate the positions of pattern matches in a string. Also `str_locate_all`.
`str_locate(fruit, "a")`

Subset Strings



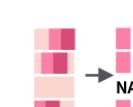
`str_sub(string, start = 1L, end = -1L)` Extract substrings from a character vector.
`str_sub(fruit, 1, 3); str_sub(fruit, -2)`



`str_subset(string, pattern)` Return only the strings that contain a pattern match.
`str_subset(fruit, "b")`



`str_extract(string, pattern)` Return the first pattern match found in each string, as a vector. Also `str_extract_all` to return every pattern match. `str_extract(fruit, "[aeiou]")`

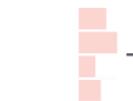


`str_match(string, pattern)` Return the first pattern match found in each string, as a matrix with a column for each () group in pattern. Also `str_match_all`.
`str_match(sentences, "(a|the) ([^]+)")`

Manage Lengths



`str_length(string)` The width of strings (i.e. number of code points, which generally equals the number of characters). `str_length(fruit)`



`str_pad(string, width, side = c("left", "right", "both"), pad = " ")` Pad strings to constant width. `str_pad(fruit, 17)`

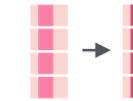


`str_trunc(string, width, side = c("right", "left", "center"), ellipsis = "...")` Truncate the width of strings, replacing content with ellipsis.
`str_trunc(fruit, 3)`

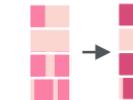


`str_trim(string, side = c("both", "left", "right"))` Trim whitespace from the start and/or end of a string. `str_trim(fruit)`

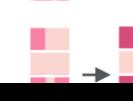
Mutate Strings



`str_sub() <- value`. Replace substrings by identifying the substrings with `str_sub()` and assigning into the results.
`str_sub(fruit, 1, 3) <- "str"`



`str_replace(string, pattern, replacement)` Replace the first matched pattern in each string. `str_replace(fruit, "a", "-")`



`str_replace_all(string, pattern, replacement)` Replace all matched patterns

Join and Split



`str_c(..., sep = "", collapse = NULL)` Join multiple strings into a single string.
`str_c(letters, LETTERS)`



`str_c(..., sep = "", collapse = "")` Collapse a vector of strings into a single string.
`str_c(letters, collapse = "")`



`str_dup(string, times)` Repeat strings times times. `str_dup(fruit, times = 2)`

Order Strings



`str_order(x, decreasing = FALSE, na_last = TRUE, locale = "en", numeric = FALSE, ...)`¹ Return the vector of indexes that sorts a character vector. `x[str_order(x)]`



`str_sort(x, decreasing = FALSE, na_last = TRUE, locale = "en", numeric = FALSE, ...)`¹ Sort a character vector.
`str_sort(x)`

Helpers

stringr Cheat Sheet

Please Take out

(character) strings

Anything surrounded by quotes(") or single quotes(').

```
> "one"  
> "1"  
> "one's"  
> ' "Hello World" '  
> "foo  
+  
+  
+ oops. I'm stuck in a string."
```



Warm Up

Type into the chat:

Are boys names or girls names more likely to end in a vowel?



babynames

year	sex	name	n	prop
			<int>	<dbl>
1880	F	Mary	7065	7.238433e-02
1880	F	Anna	2604	2.667923e-02
1880	F	Emma		
1880	F	Elizabeth		
1880	F	Minnie		
1880	F	Margaret		
1880	F	Ida		
1880	F	Alice	1414	1.448711e-02
1880	F	Bertha	1320	1.352404e-02
1880	F	Sarah	1288	1.319618e-02

How can we build the proportion of boys and girls whose name ends in a vowel?

1-10 of 1,858,689 rows

Previous

1

2

3

4

5

6

...

100 Next



Most useful skills

1. How to extract/ replace substrings
2. How to find matches for patterns
3. Regular expressions



stringr



Simple, consistent functions for working with strings.

```
# install.packages("tidyverse")
library(stringr)
```



```
install.packages("tidyverse")
```

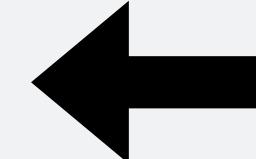
does the equivalent of

```
install.packages("ggplot2")
install.packages("dplyr")
install.packages("tidyr")
install.packages("readr")
install.packages("purrr")
install.packages("tibble")
install.packages("stringr")
install.packages("hms")
install.packages("lubridate")
install.packages("stringr")
install.packages("forcats")
install.packages("DBI")
install.packages("haven")
install.packages("httr")
install.packages("jsonlite")
install.packages("readxl")
install.packages("rvest")
install.packages("xml2")
install.packages("modelr")
install.packages("broom")
```

```
library("tidyverse")
```

does the equivalent of

```
library("ggplot2")
library("dplyr")
library("tidyr")
library("readr")
library("purrr")
library("tibble")
library("stringr")
```



str_sub()

Extract or replace portions of a string with **str_sub()**

```
str_sub(string, start = 1, end = -1)
```

**string(s) to
manipulate**

**position of first
character to extract
within each string**

**position of last
character to extract
within each string**



Quiz

What will this return?

```
str_sub("Garrett", 1, 2)
```

Quiz

What will this return?

```
str_sub("Garrett", 1, 2)
```

"Ga"

Quiz

What will this return?

```
str_sub("Garrett", 1, 1)
```

Quiz

What will this return?

```
str_sub("Garrett", 1, 1)
```

"G"

Quiz

What will this return?

```
str_sub("Garrett", 2)
```

Quiz

What will this return?

```
str_sub("Garrett", 2)
```

"arrett"

Quiz

What will this return?

```
str_sub("Garrett", -3)
```

Quiz

What will this return?

```
str_sub("Garrett", -3)
```

"ett"

Quiz

What will this return?

```
g <- "Garrett"
```

```
str_sub(g, -3) <- "eth"
```

```
g
```

Quiz

What will this return?

```
g <- "Garrett"
```

```
str_sub(g, -3) <- "eth"
```

```
g
```

"Garreth"

Your Turn 2

In your group, fill in the blanks to:

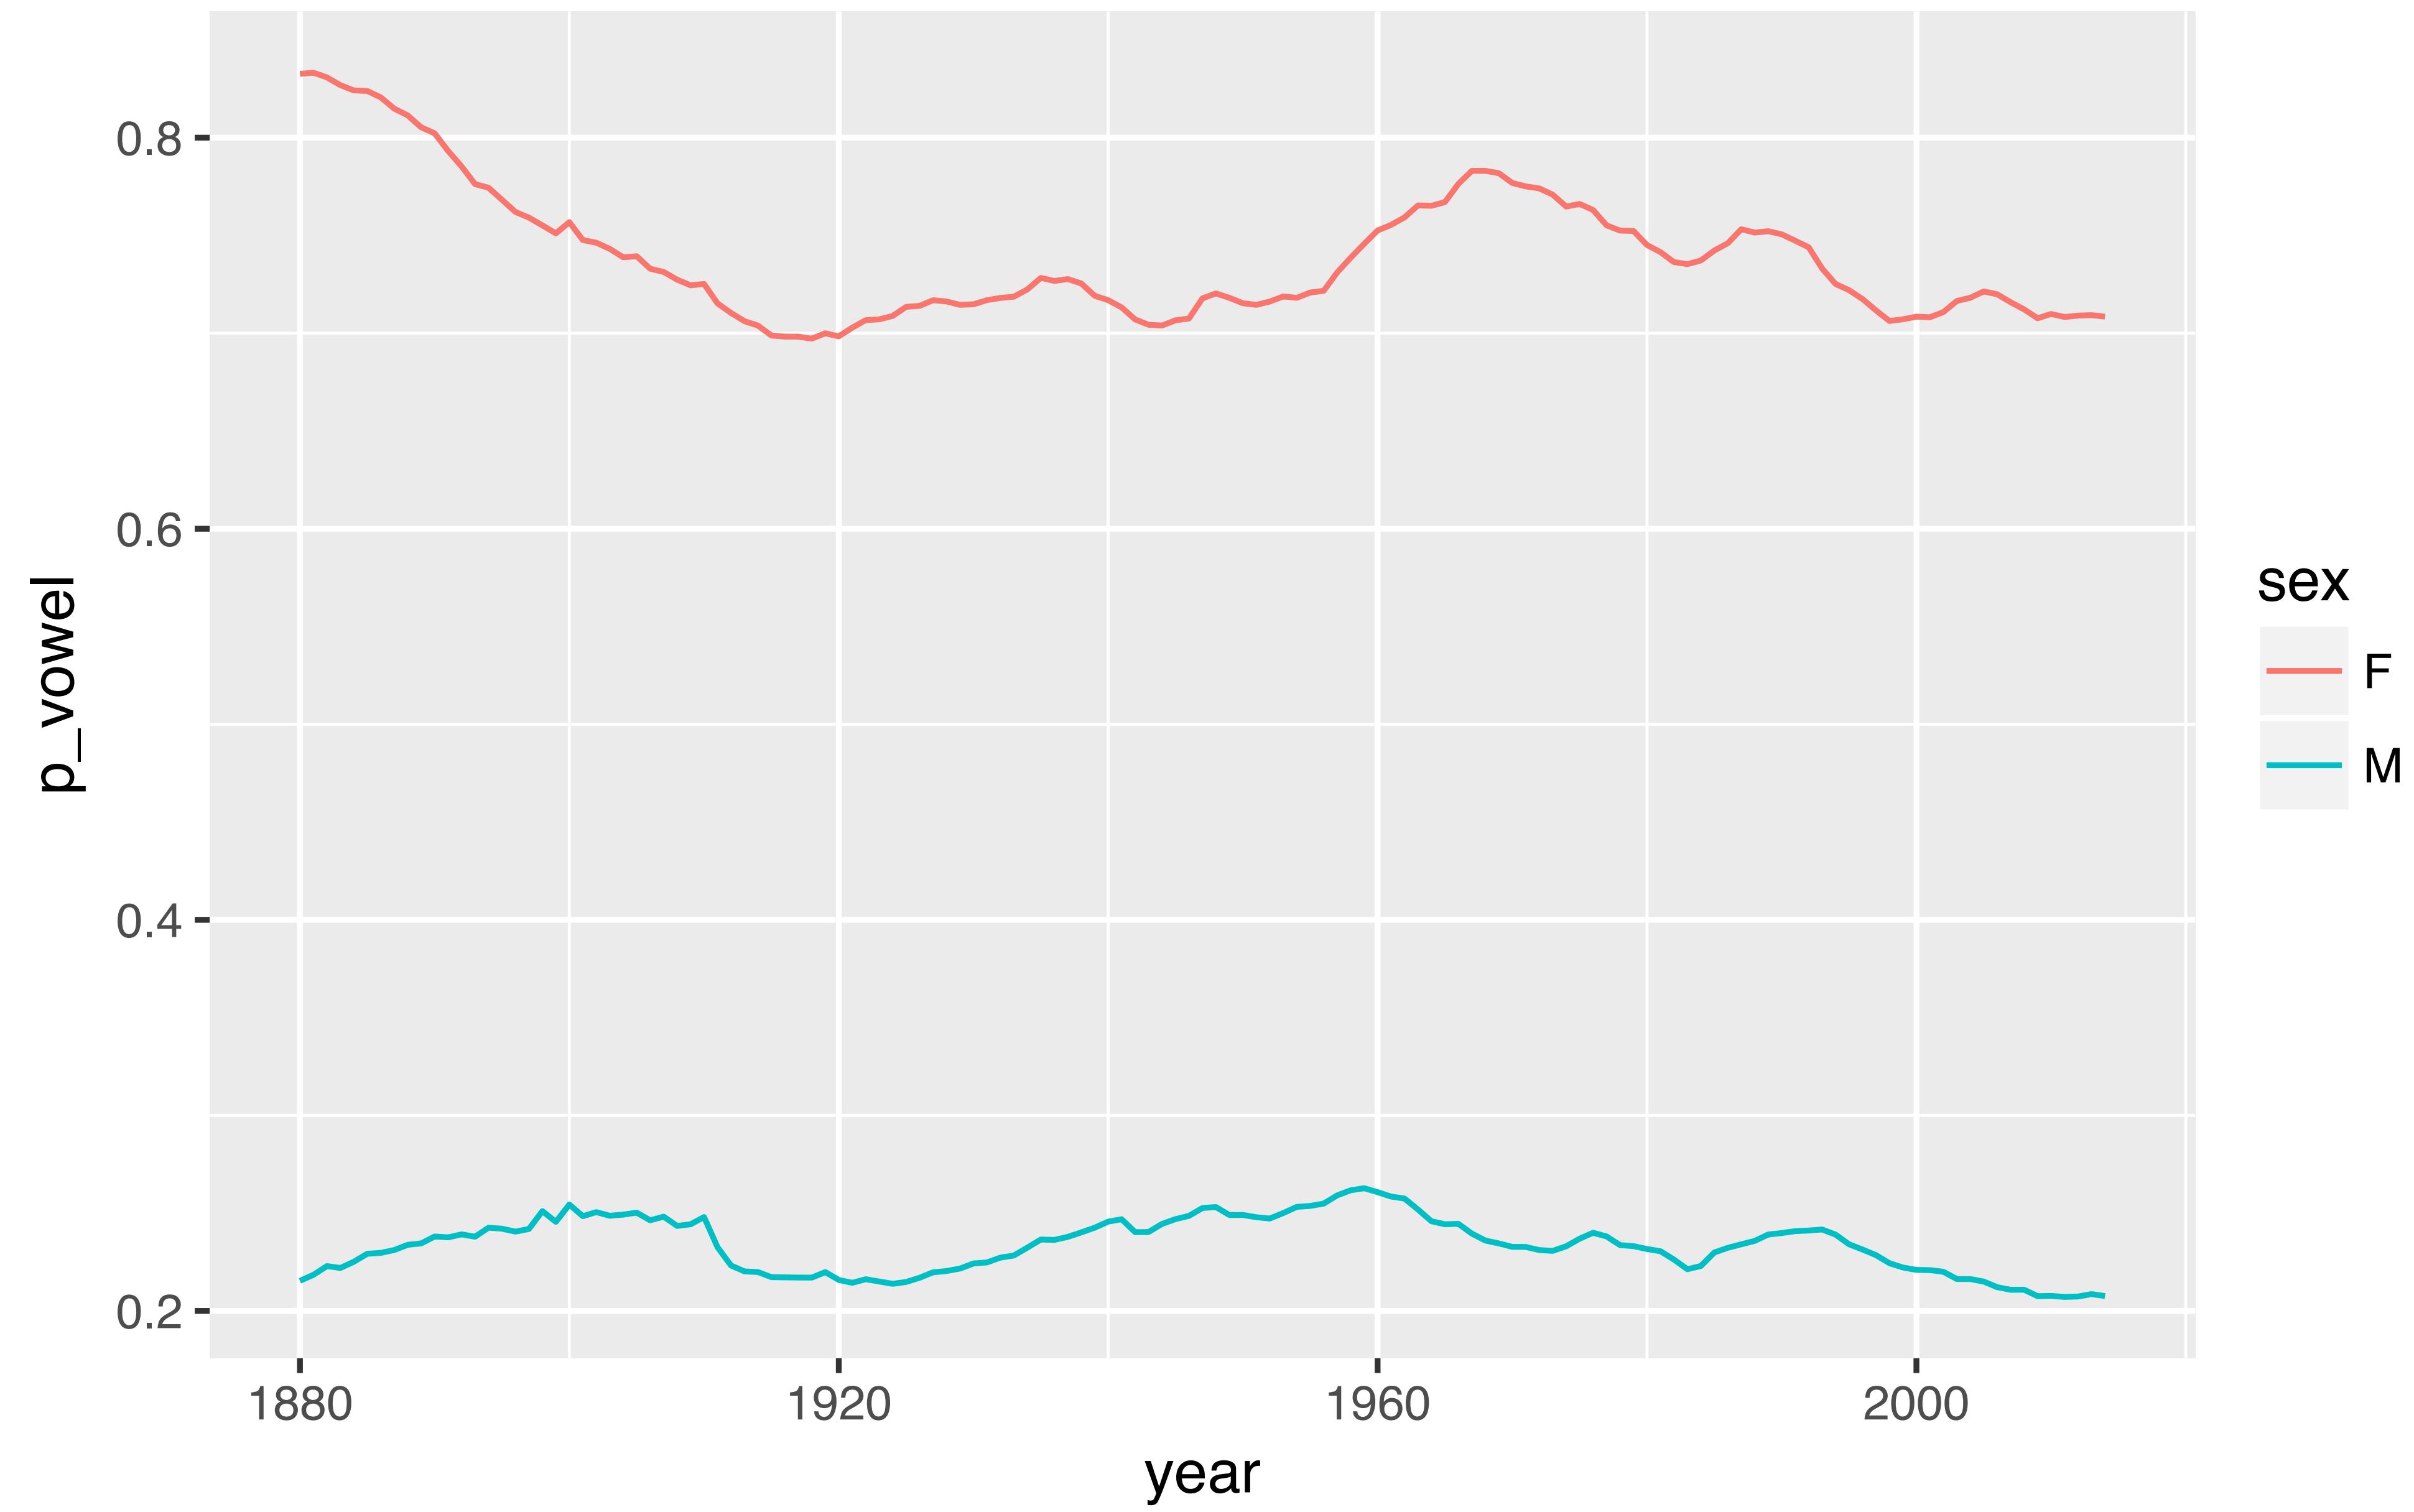
1. Isolate the last letter of every name
2. and create a logical variable that displays whether the last letter is one of "a", "e", "i", "o", "u", or "y".
3. Use a **weighted mean** to calculate the proportion of children whose name ends in a vowel (by year and sex).
4. and then display the results as a line plot.



```
babynames %>%  
  mutate(last = str_sub(name, -1),  
         vowel = last %in% c("a", "e", "i", "o", "u", "y")) %>%  
  group_by(year, sex) %>%  
  summarise(p_vowel = weighted.mean(vowel, n)) %>%  
  ggplot() +  
    geom_line(mapping = aes(year, p_vowel, color = sex))
```

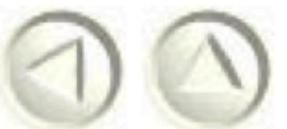


Proportion of names that end in a vowel



```
help(package = stringr)
```

Simple, Consistent Wrappers for Common String Operations



Documentation for package ‘stringr’ version 1.2.0

- [DESCRIPTION file](#).
- [User guides, package vignettes and other documentation](#).

Help Pages

[boundary](#)

Control matching behaviour with modifier functions.

[case](#)

Convert case of a string.

[coll](#)

Control matching behaviour with modifier functions.

[fixed](#)

Control matching behaviour with modifier functions.

[fruit](#)

Sample character vectors for practicing string manipulations.

[invert_match](#)

Switch location of matches to location of non-matches.

[modifiers](#)

Control matching behaviour with modifier functions.

[regex](#)

Control matching behaviour with modifier functions.

[sentences](#)

Sample character vectors for practicing string manipulations.



Description

Vectorised over string and pattern. Equivalent to `grepl(pattern, x)`. See [str\(which\(\)\)](#) for an equivalent to `grep(pattern, x)`.

Usage

```
str_detect(string, pattern, negate = FALSE)
```

Arguments

`string` Input vector. Either a character vector, or something coercible to one.

`pattern` Pattern to look for.

The default interpretation is a regular expression, as described in [stringi::stringi-search-regex](#). Control

Regular Expressions

Advanced, flexible string detection

Need to Know

Pattern arguments in stringr are interpreted as regular expressions *after* any special characters have been parsed.

In R, you write regular expressions as *strings*, sequences of characters surrounded by quotes ("") or single quotes('').

Some characters cannot be represented directly in an R String. These must be represented as **special characters**, sequences of characters that have a specific meaning., e.g.

Special Character	Represents
\\\	\
\"	"
\n	new line

Run ?"" to see a complete list

Because of this, whenever a \ appears in a regular expression, you must write it as \\ in the string that represents the regular expression.

Use `writeLines()` to see how R views your string after all special characters have been parsed.

```
writeLines("\\".)  
# .  
  
writeLines("\\ is a backslash")  
# \ is a backslash
```

INTERPRETATION

Patterns in stringr are interpreted as regexes To change this default, wrap the pattern in one of:

```
regex(pattern, ignore_case = FALSE, multiline =  
      FALSE, comments = FALSE, dotall = FALSE, ...)  
Modifies a regex to ignore cases, match end of  
lines as well of end of strings, allow R comments  
within regex's, and/or to have . match everything  
including \n.  
str_detect("I", regex("i", TRUE))
```

Regular Expressions -

Regular expressions, or *regexp*s, are a concise language for describing patterns in strings.

MATCH CHARACTERS

string (type this)	regexp (to mean this)	matches (which matches this)
a (etc.)	a (etc.)	a (etc.)
\\.	\\. (dot)	.
\\!	\\! (exclamation mark)	!
\\?	\\? (question mark)	?
\\\\	\\\\ (backslash)	\\
\\(\\((left parenthesis)	(
\\)	\\) (right parenthesis))
\\{	\\{ (left brace)	{
\\}	\\} (right brace)	}
\\n	\\n (newline)	new line (return)
\\t	\\t (tab)	tab
\\s	\\s (any whitespace)	any whitespace (\\$ for non-whitespaces)
\\d	\\d (any digit)	any digit (\\$ for non-digits)
\\w	\\w (any word character)	any word character (\\$W for non-word chars)
\\b	\\b (word boundaries)	word boundaries
[:digit:] ¹	[:digit:]	digits
[:alpha:] ¹	[:alpha:]	letters
[:lower:] ¹	[:lower:]	lowercase letters
[:upper:] ¹	[:upper:]	uppercase letters
[:alnum:] ¹	[:alnum:]	letters and numbers
[:punct:] ¹	[:punct:]	punctuation
[:graph:] ¹	[:graph:]	letters, numbers, and punctuation
[:space:] ¹	[:space:]	space characters (i.e. \s)
[:blank:] ¹	[:blank:]	space and tab (but not new line)
.	.	every character except a new line

see <- function(rx) str_view_all("abc ABC 123\\t.!?\n", rx)

example

see("a")	abc ABC 123 .!?\n
see("\\.")	abc ABC 123 .!?\n
see("\\!")	abc ABC 123 .!?\n
see("\\?")	abc ABC 123 .!?\n
see("\\\\")	abc ABC 123 .!?\n
see("\\()")	abc ABC 123 .!?\n
see("\\)")	abc ABC 123 .!?\n
see("\\{")	abc ABC 123 .!?\n
see("\\}")	abc ABC 123 .!?\n
see("\\n")	abc ABC 123 .!?\n
see("\\t")	abc ABC 123 .!?\n
see("\\s")	abc ABC 123 .!?\n
see("\\d")	abc ABC 123 .!?\n
see("\\w")	abc ABC 123 .!?\n
see("\\b")	abc ABC 123 .!?\n
see("[[:digit:]]")	abc ABC 123 .!?\n
see("[[:alpha:]]")	abc ABC 123 .!?\n
see("[[:lower:]]")	abc ABC 123 .!?\n
see("[[:upper:]]")	abc ABC 123 .!?\n
see("[[:alnum:]]")	abc ABC 123 .!?\n
see("[[:punct:]]")	abc ABC 123 .!?\n
see("[[:graph:]]")	abc ABC 123 .!?\n
see("[[:space:]]")	abc ABC 123 .!?\n
see("[[:blank:]]")	abc ABC 123 .!?\n
see("")	abc ABC 123 .!?\n

¹ Many base R functions require classes to be wrapped in a second set of [], e.g. [[:digit:]]

ALTERNATES

alt <- function(rx) str_view_all("abcde", rx)

regexp	matches	example
ab d	or	alt("ab d") abcde
[abe]	one of	alt("[abe]") abcde
[^abe]	anything but	alt("[^abe]") abcde
[a-c]	range	alt("[a-c]") abcde

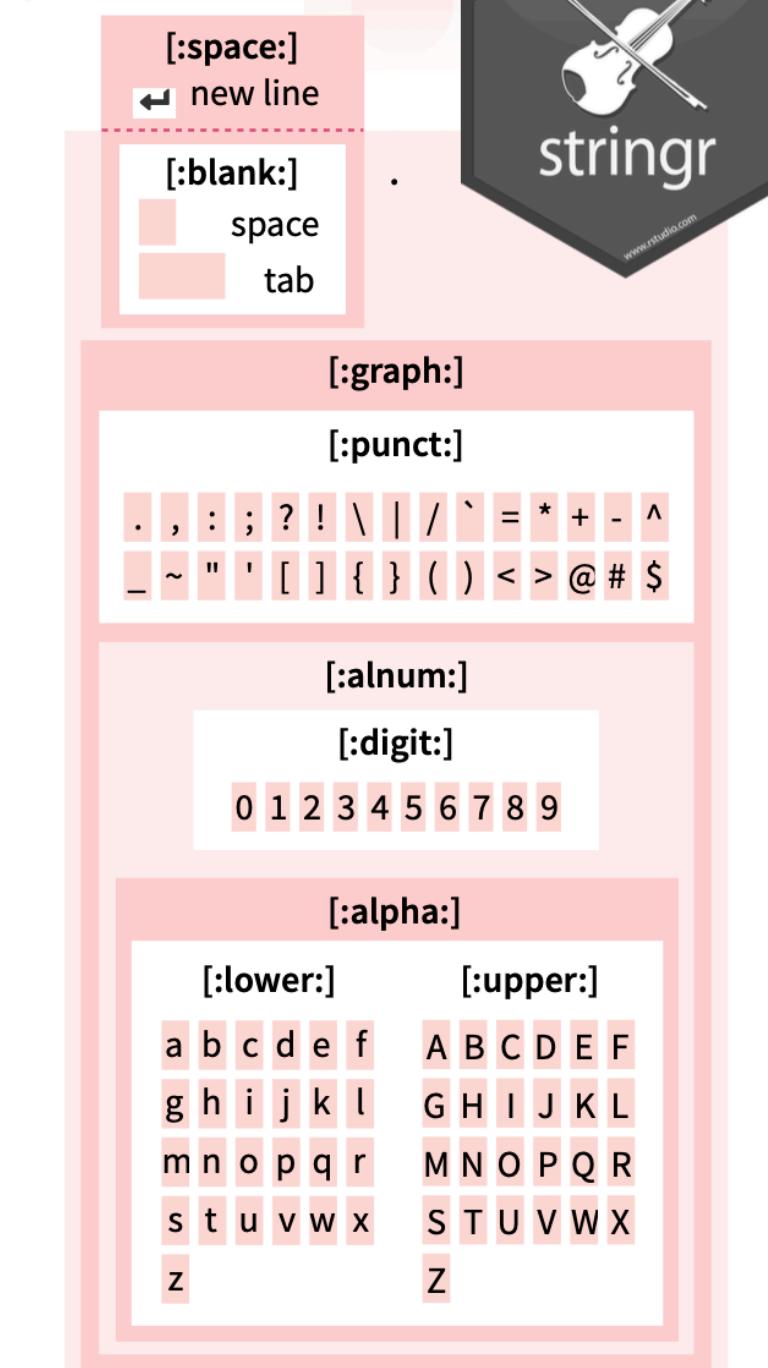
QUANTIFIERS

quant <- function(rx) str_view_all(".a.aa.aaa", rx)

regexp	matches	example
a?	zero or one	quant("a?") .a.aa.aaa
a*	zero or more	quant("a*") .a.aa.aaa
a+	one or more	quant("a+") .a.aa.aaa
a{n}	exactly n	quant("a{2}") .a.aa.aaa

quant <- function(rx) str_view_all(".a.aa.aaa", rx)

regexp	matches	example
a?	zero or one	quant("a?") .a.aa.aaa
a*	zero or more	quant("a*") .a.aa.aaa
a+	one or more	quant("a+") .a.aa.aaa
a{n}	exactly n	quant("a{2}") .a.aa.aaa



Regular Expressions

Advanced, flexible string detection
Page 2 of Cheat Sheet

WHENEVER I LEARN A
NEW SKILL I CONCOCT
ELABORATE FANTASY
SCENARIOS WHERE IT
LETS ME SAVE THE DAY.

OH NO! THE KILLER
MUST HAVE FOLLOWED
HER ON VACATION!



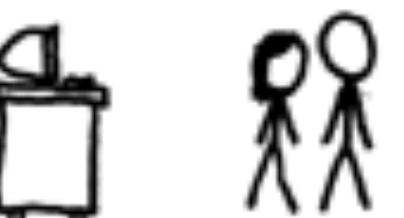
BUT TO FIND THEM WE'D HAVE TO SEARCH
THROUGH 200 MB OF EMAILS LOOKING FOR
SOMETHING FORMATTED LIKE AN ADDRESS!



EVERYBODY STAND BACK.



I KNOW REGULAR
EXPRESSIONS.



\w*[aeiouy]\b

Symbol	Meaning
\w	"word character"
*	"any number of"
[aeiouy]	"any character in between the []"
\b	"word boundary"

```
> str_detect("Ari", "\\w*[aeiouy]\\b")
[1] TRUE
> str_detect("Bob", "\\w*[aeiouy]\\b")
[1] FALSE
```

\ becomes \\

hello@example.com

Create a RegEx for an email address...
How hard can it be?

The [fully RFC 822 compliant regex](#) is inefficient and obscure because of its length. Fortunately, RFC 822 was superseded twice and the current specification for email addresses is [RFC 5322](#).

2553 RFC 5322 leads to a regex that can be understood if studied for a few minutes and is efficient enough for actual use.

✓ One RFC 5322 compliant regex can be found at the top of the page at <http://emailregex.com/> but uses the IP address pattern that is floating around the internet with a bug that allows `00` for any of the unsigned byte decimal values in a dot-delimited address, which is illegal. The rest of it appears to be consistent with the RFC 5322 grammar and passes several tests using `grep -Po`, including cases domain names, IP addresses, bad ones, and account names with and without quotes.

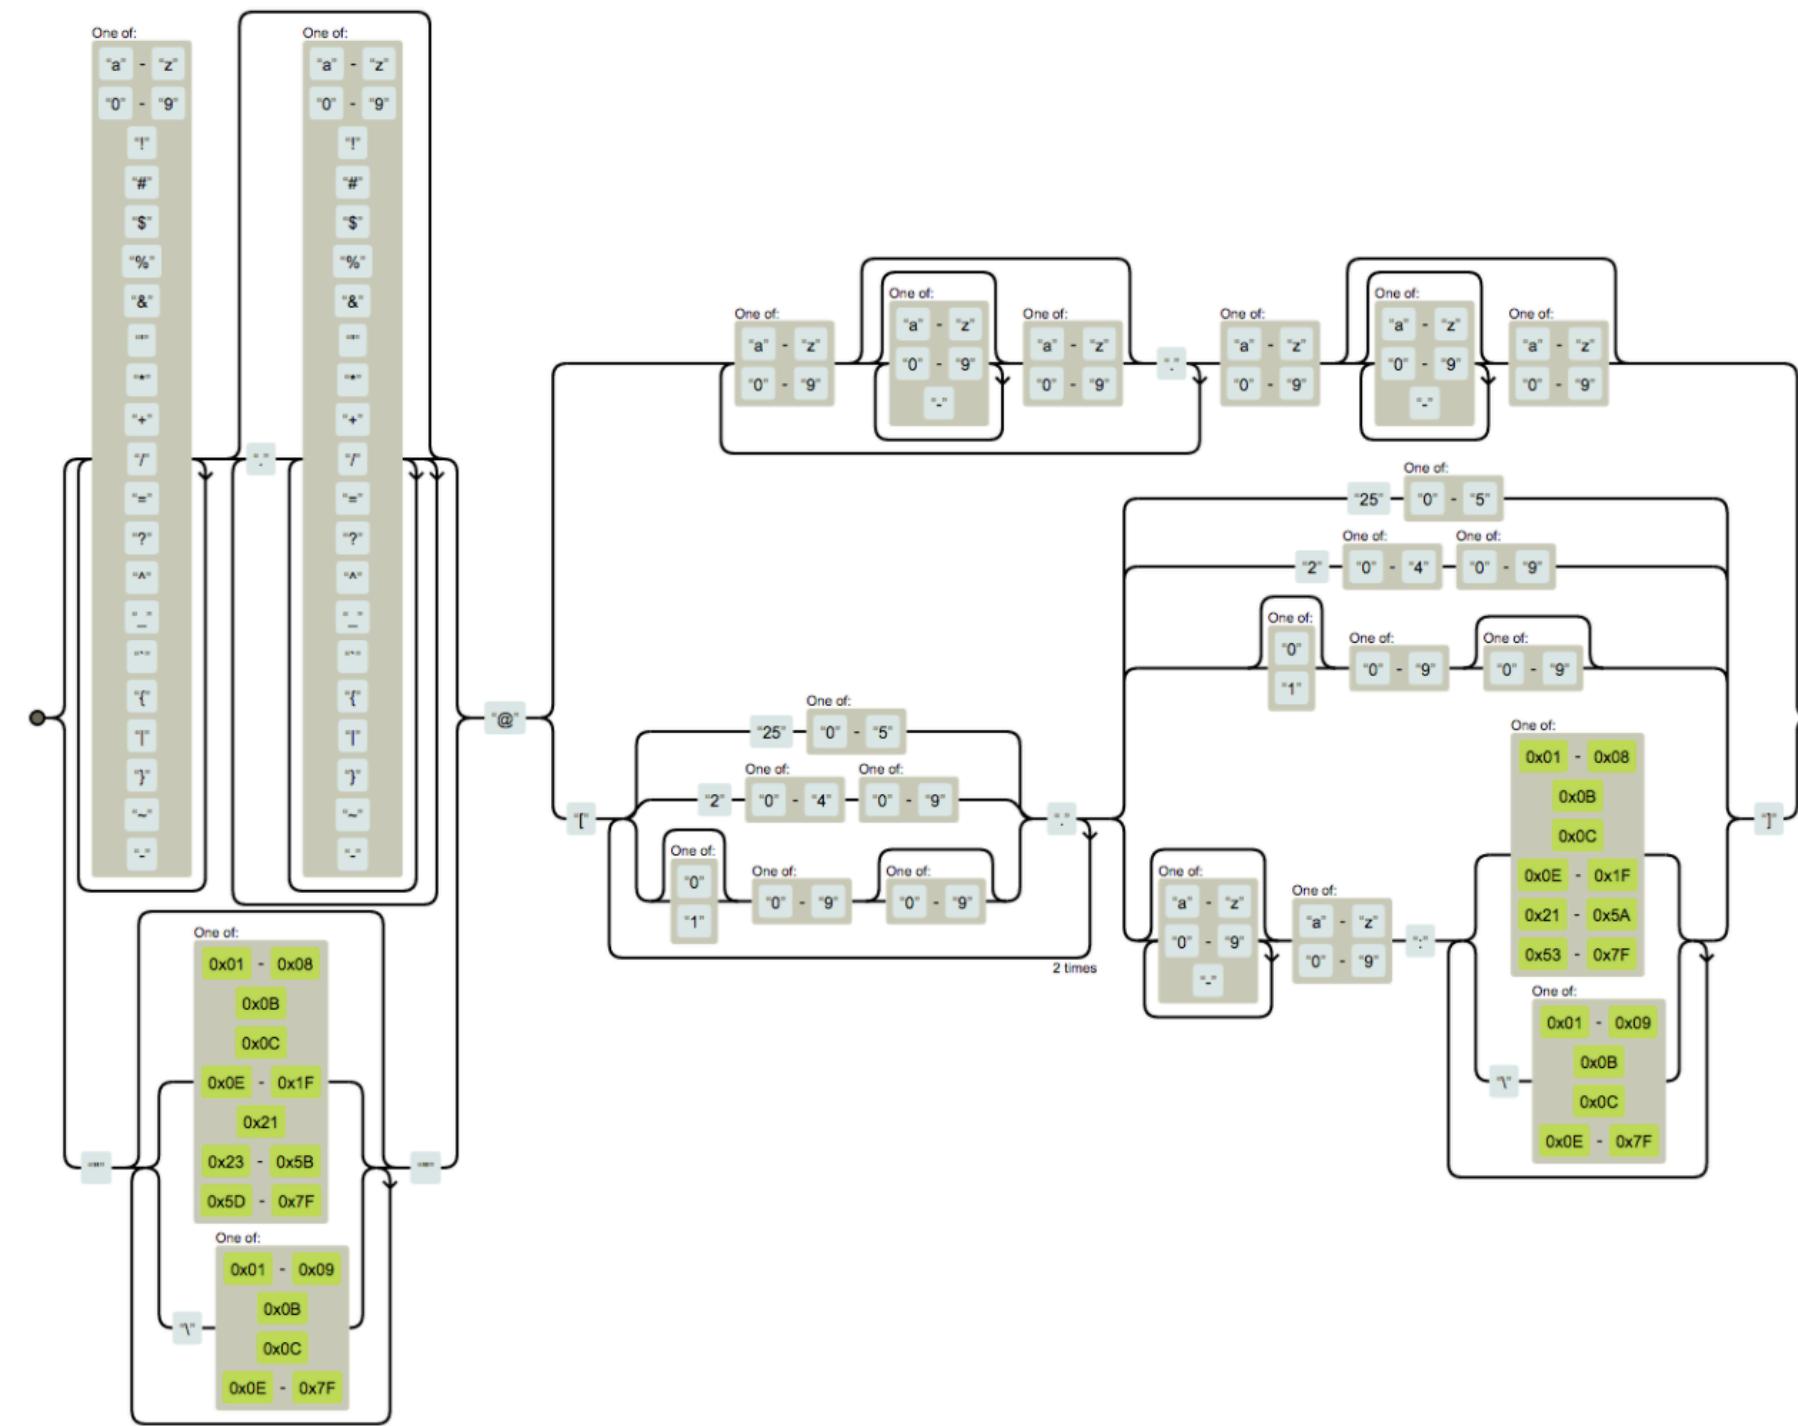
Correcting the `00` bug in the IP pattern, we obtain a working and fairly fast regex. (Scrape the rendered version, not the markdown, for actual code.)

```
(?:[a-z0-9!#$%&'*+/=?^_`{|}~-]+(?:\.[a-z0-9!#$%&'*+/=?^_`{|}~-]+)*|"(?:[x01-x08\x0b\x0c\x0e-\x1f\x21\x23-\x5b\x5d-\x7f]|\\[\x01-\x09\x0b\x0c\x0e-\x7f])*")@(?:(?:[a-z0-9](?:[a-z0-9]*[a-z0-9])?|.)+[a-z0-9](?:[a-z0-9]*[a-z0-9])?|[(?:(?:2(5[0-5][0-4][0-9])|1[0-9][0-9])[1-9]?[0-9]))|.]{3}(?:2(5[0-5][0-4][0-9])|1[0-9][0-9])[1-9]?[0-9])|[a-z0-9]*[a-z0-9](?:[\x01-\x08\x0b\x0c\x0e-\x1f\x21-\x5a\x53-\x7f]|\\[\x01-\x09\x0b\x0c\x0e-\x7f])+))]
```

or:

Not Human Readable

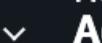
Here is [diagram of finite state machine](#) for above regexp which is more clear than regexp itself



The more sophisticated patterns in Perl and PCRE (regex library used e.g. in PHP) can [correctly parse RFC 5322 without a hitch](#). Python and C# can do that too, but they use a different syntax from those first two. However, if you are forced to use one of the many less powerful pattern-matching languages, then it's best to use a real parser.

"RegEx Visualizer"

"Railroad Diagrams"

amazon  All regular expressions   Hello, Sign in Account Returns & Orders 

All Holiday Deals Gift Cards Best Sellers Prime Customer Service Find a Gift New Releases Whole Foods Early Black Friday deals

1-48 of 306 results for "regular expressions" Sort by: Featured

Eligible for Free Shipping

Free Shipping by Amazon
All customers get FREE Shipping on orders over \$25 shipped by Amazon

Kindle Unlimited

Kindle Unlimited Eligible

Department

Books

Programming Languages
Software Design, Testing & Engineering
Microsoft Programming
Computer Science
MySQL Guides
Web Development & Design
Programming
Other Databases

Kindle Store

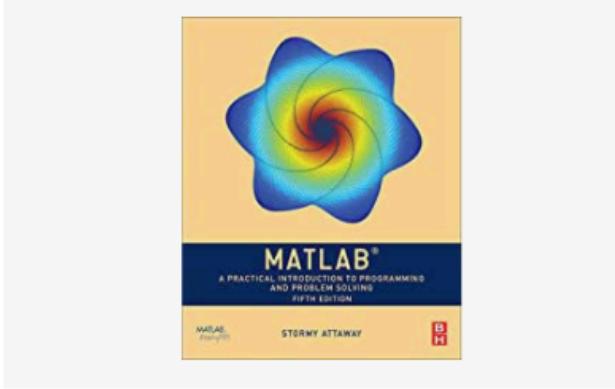
Computer Programming
[See All 3 Departments](#)

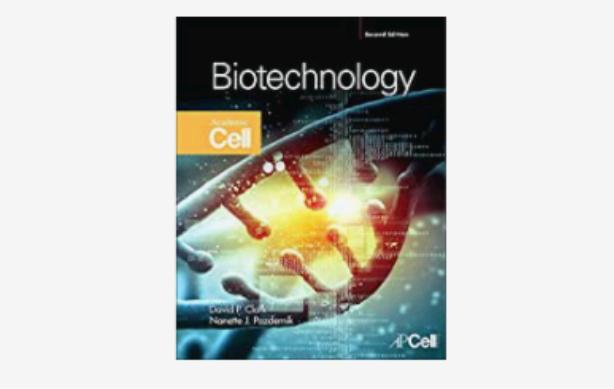
Avg. Customer Review

 & Up
 & Up

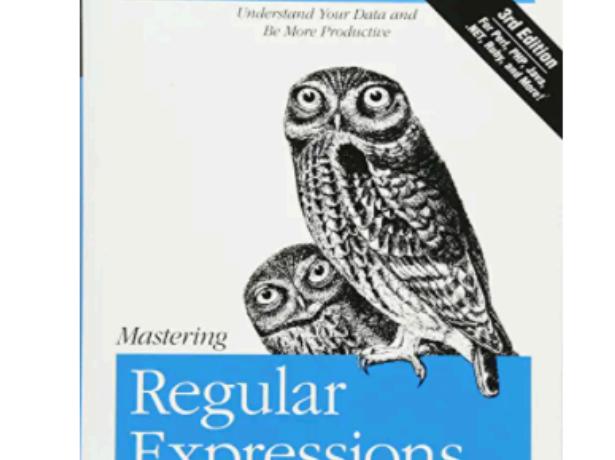
 **Empowering your research in Engineering**
[Shop Elsevier Science & Technology >](#)

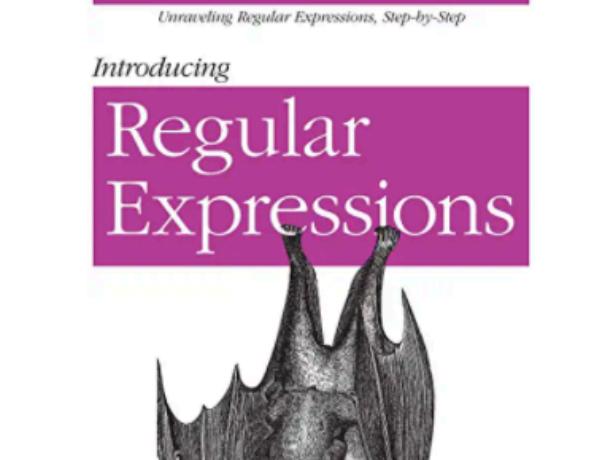

Analyzing Social Media Networks with NodeXL: Insights from a Connected...


MATLAB: A Practical Introduction to Programming and Problem Solving
4.5 stars 73 reviews
prime


Biotechnology
4.5 stars 26 reviews
prime


Mastering Python Regular Expressions


Mastering Regular Expressions
Unraveling Regular Expressions, Step-by-Step


Introducing Regular Expressions

Sponsored

Tons of books on Reg Ex

Online tutorials and cheat sheets as well

Factors

R

Factors withforcats :: CHEAT SHEET

The **forcats** package provides tools for working with factors, which are R's data structure for categorical data.

Factors

R represents categorical data with factors. A **factor** is an integer vector with a **levels** attribute that stores a set of mappings between integers and categorical values. When you view a factor, R displays not the integers, but the values associated with them.

Create a factor with factor()
`factor(x = character(), levels, labels = levels, exclude = NA, ordered = is.ordered(x), nmax = NA)` Convert a vector to a factor. Also **as_factor()**.
`f <- factor(c("a", "c", "b", "a"), levels = c("a", "b", "c"))`

Return its levels with levels()
`levels(x)` Return/set the levels of a factor. `levels(f); levels(f) <- c("x", "y", "z")`

Use unclass() to see its structure

Inspect Factors

fct_count(f, sort = FALSE) Count the number of values with each level. `fct_count(f)`

fct_unique(f) Return the unique values, removing duplicates. `fct_unique(f)`

Change the order of levels

fct_relevel(f, ..., after = 0L) Manually reorder factor levels.
`fct_relevel(f, c("b", "c", "a"))`

fct_infreq(f, ordered = NA) Reorder levels by the frequency in which they appear in the data (highest frequency first).
`f3 <- factor(c("c", "c", "a"))`
`fct_infreq(f3)`

fct_inorder(f, ordered = NA) Reorder levels by order in which they appear in the data.
`fct_inorder(f2)`

fct_rev(f) Reverse level order.
`f4 <- factor(c("a", "b", "c"))`
`fct_rev(f4)`

fct_shift(f) Shift levels to left or right, wrapping around end.
`fct_shift(f4)`

fct_shuffle(f, n = 1L) Randomly permute order of factor levels.
`fct_shuffle(f4)`

Change the value of levels

fct_recode(f, ...) Manually change levels. Also **fct_relabel** which obeys purrr::map syntax to apply a function or expression to each level.
`fct_recode(f, v = "a", x = "b", z = "c")`
`fct_relabel(f, ~ paste0("x", x))`

fct_anon(f, prefix = "") Anonymize levels with random integers. `fct_anon(f)`

fctCollapse(f, ...) Collapse levels into manually defined groups.
`fct_collapse(f, x = c("a", "b"))`

fct_lump(f, n, prop, w = NULL, other_level = "Other", ties.method = c("min", "average", "first", "last", "random", "max")) Lump together least/most common levels into a single level. Also **fct_lump_min**.
`fct_lump(f, n = 1)`

fct_other(f, keep, drop, other_level = "Other") Replace levels with "other."
`fct_other(f, keep = c("a", "b"))`

forcats Cheat Sheet

Please take out



factors

R's representation of categorical data. Consists of:

1. A set of **values**
2. An ordered set of **valid levels**

```
eyes <- factor(x = c("blue", "green", "green"),
                 levels = c("blue", "brown", "green"))

eyes
## [1] blue green green
## Levels: blue brown green
```



factors

Stored as an integer vector with a levels attribute

```
unclass(eyes)
## 1 3 3
## attr(,"levels")
## "blue" "brown" "green"
```



forcats



Simple functions for working with factors.

```
# install.packages("tidyverse")
library(forcats)
```



Warm Up

Decide in your group:

Which religions watch the least TV?

Do married people watch more or less TV than single people?



gss_cat

```
library(forcats)  
gss_cat
```

A sample of data from the General Social Survey, a long-running US survey conducted by NORC at the University of Chicago.

tvhours	marital	age	race	partyid	relig
<int>	<fctr>	<int>	<fctr>	<fctr>	<fctr>
12	Never married	26	White	Ind,near rep	Protestant
NA	Divorced	48	White	Not str republican	Protestant
2	Widowed	67	White	Independent	Protestant
4	Never married	39	White	Ind,near rep	Orthodox-christian
1	Divorced	25	White	Not str democrat	None
NA	Married	25	White	Strong democrat	Protestant
3	Never married	36	White	Not str republican	Christian
NA	Divorced	44	White	Ind,near dem	Protestant

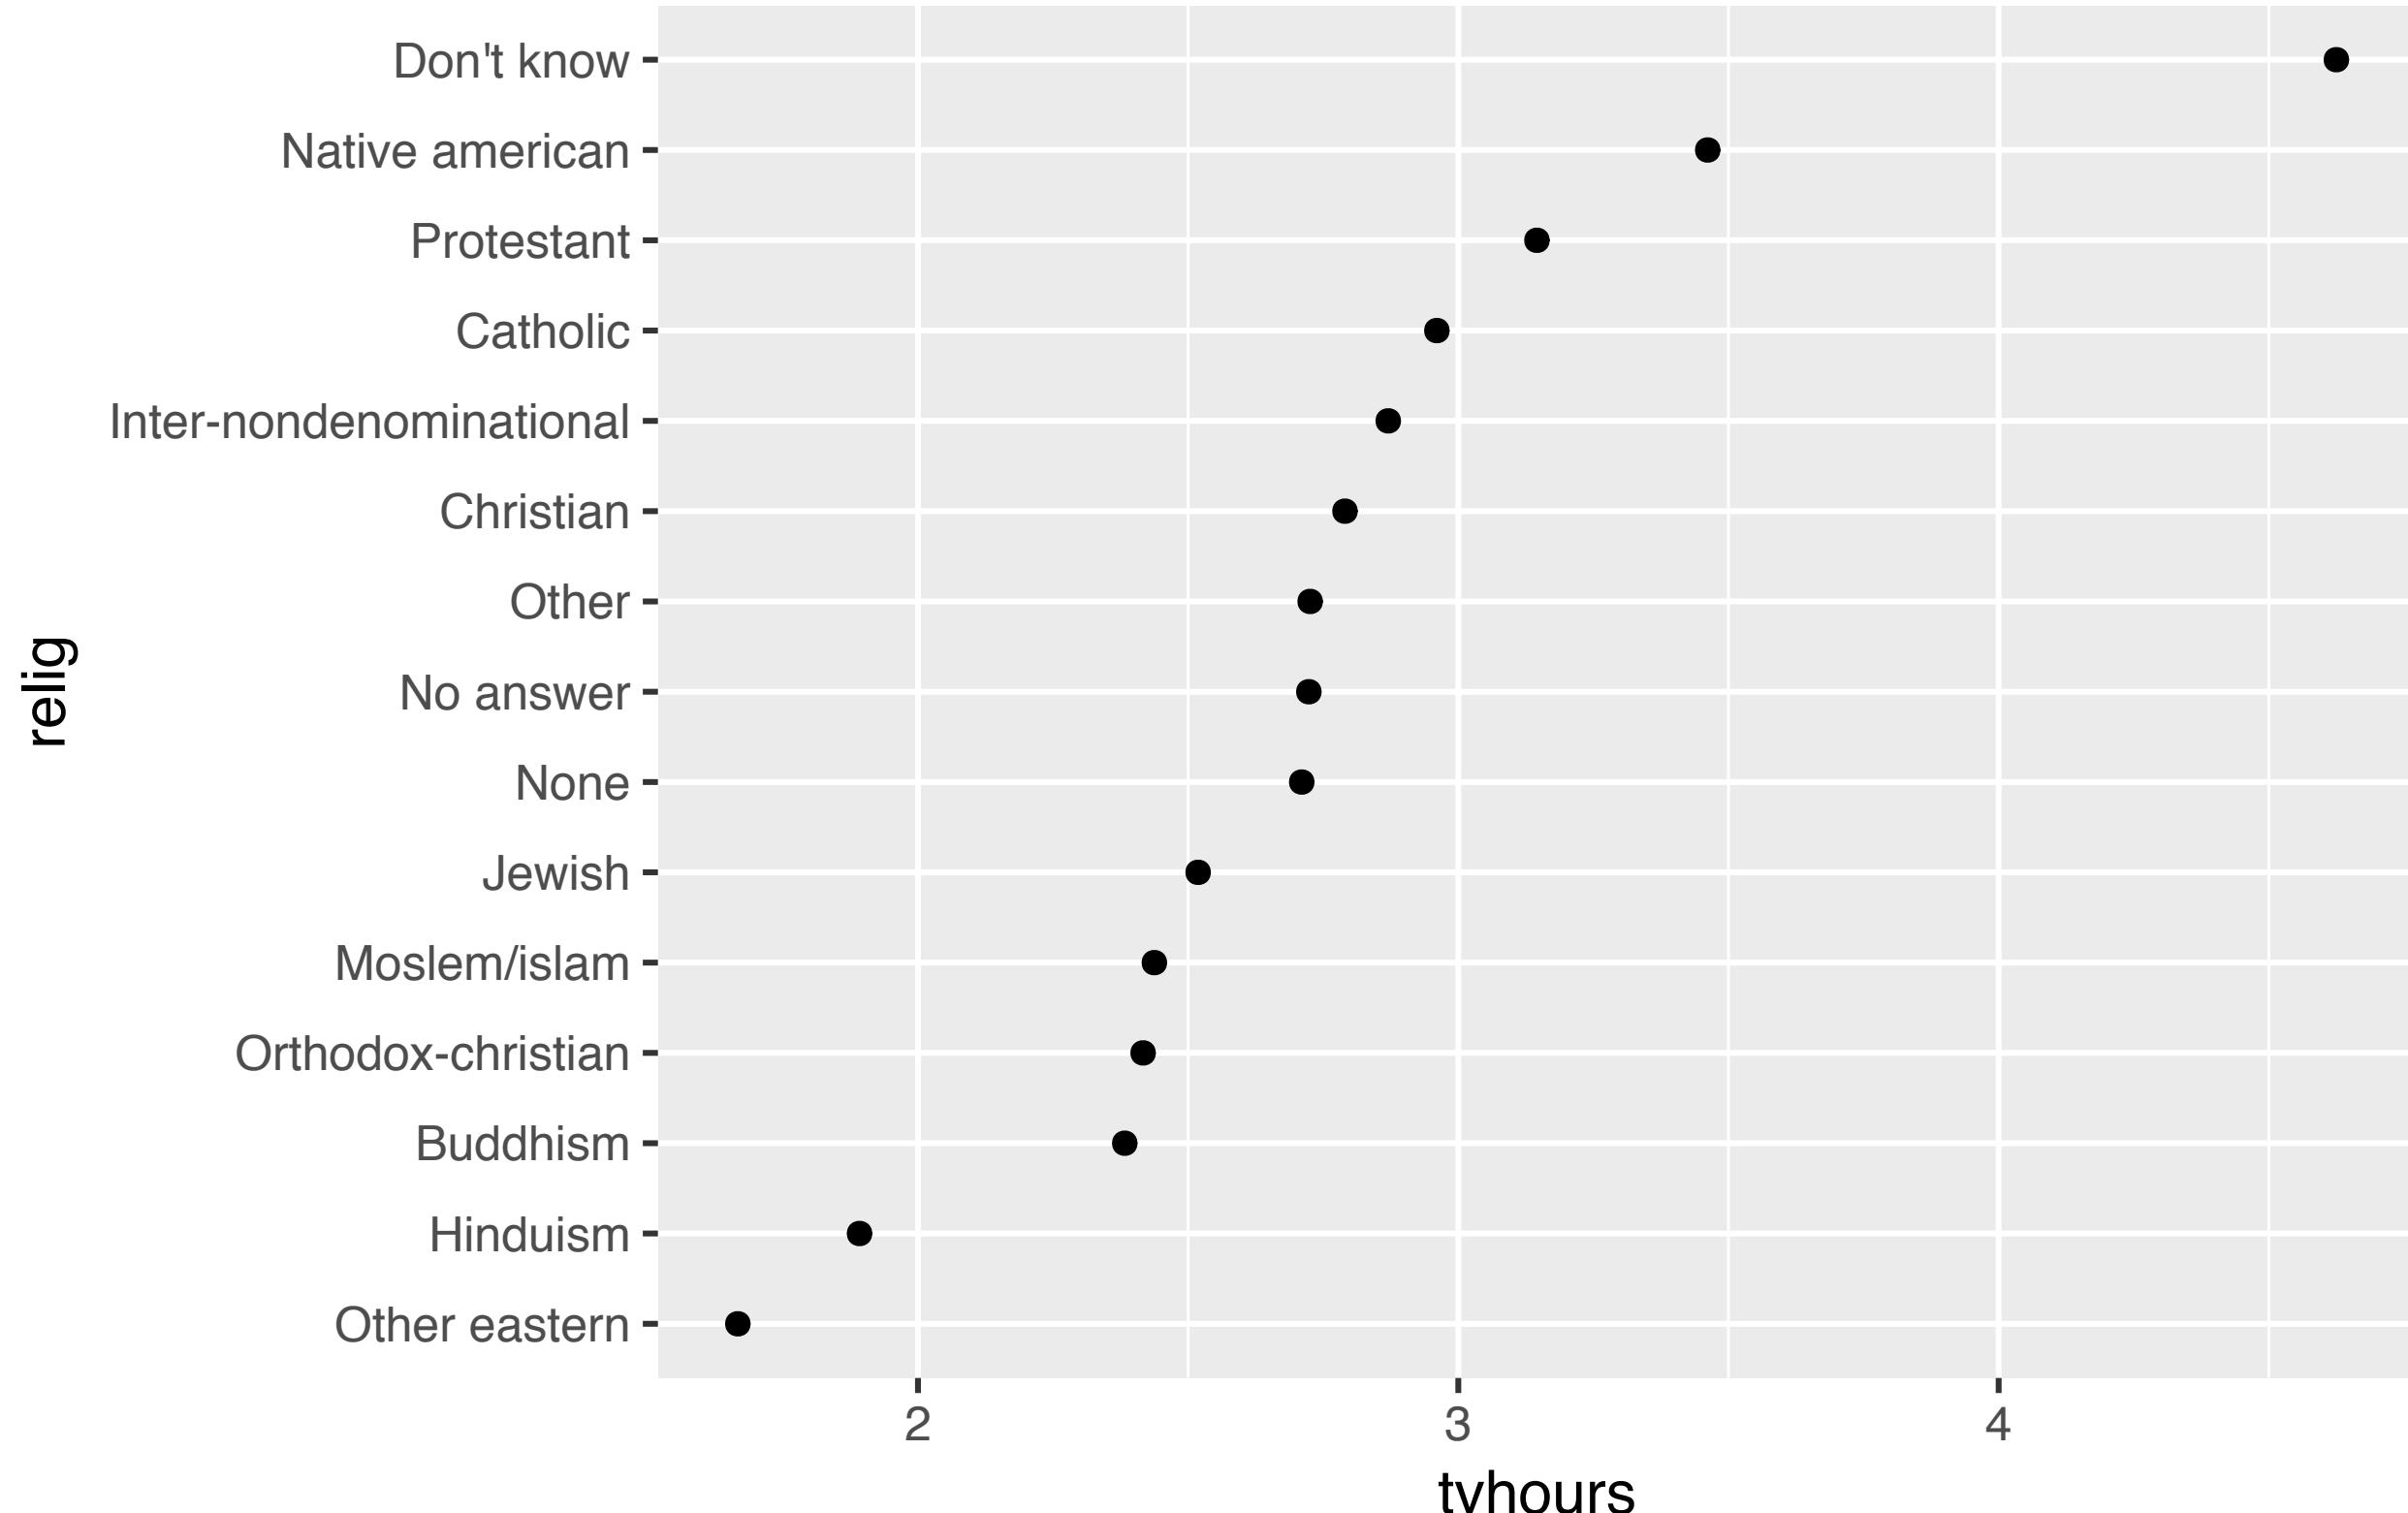
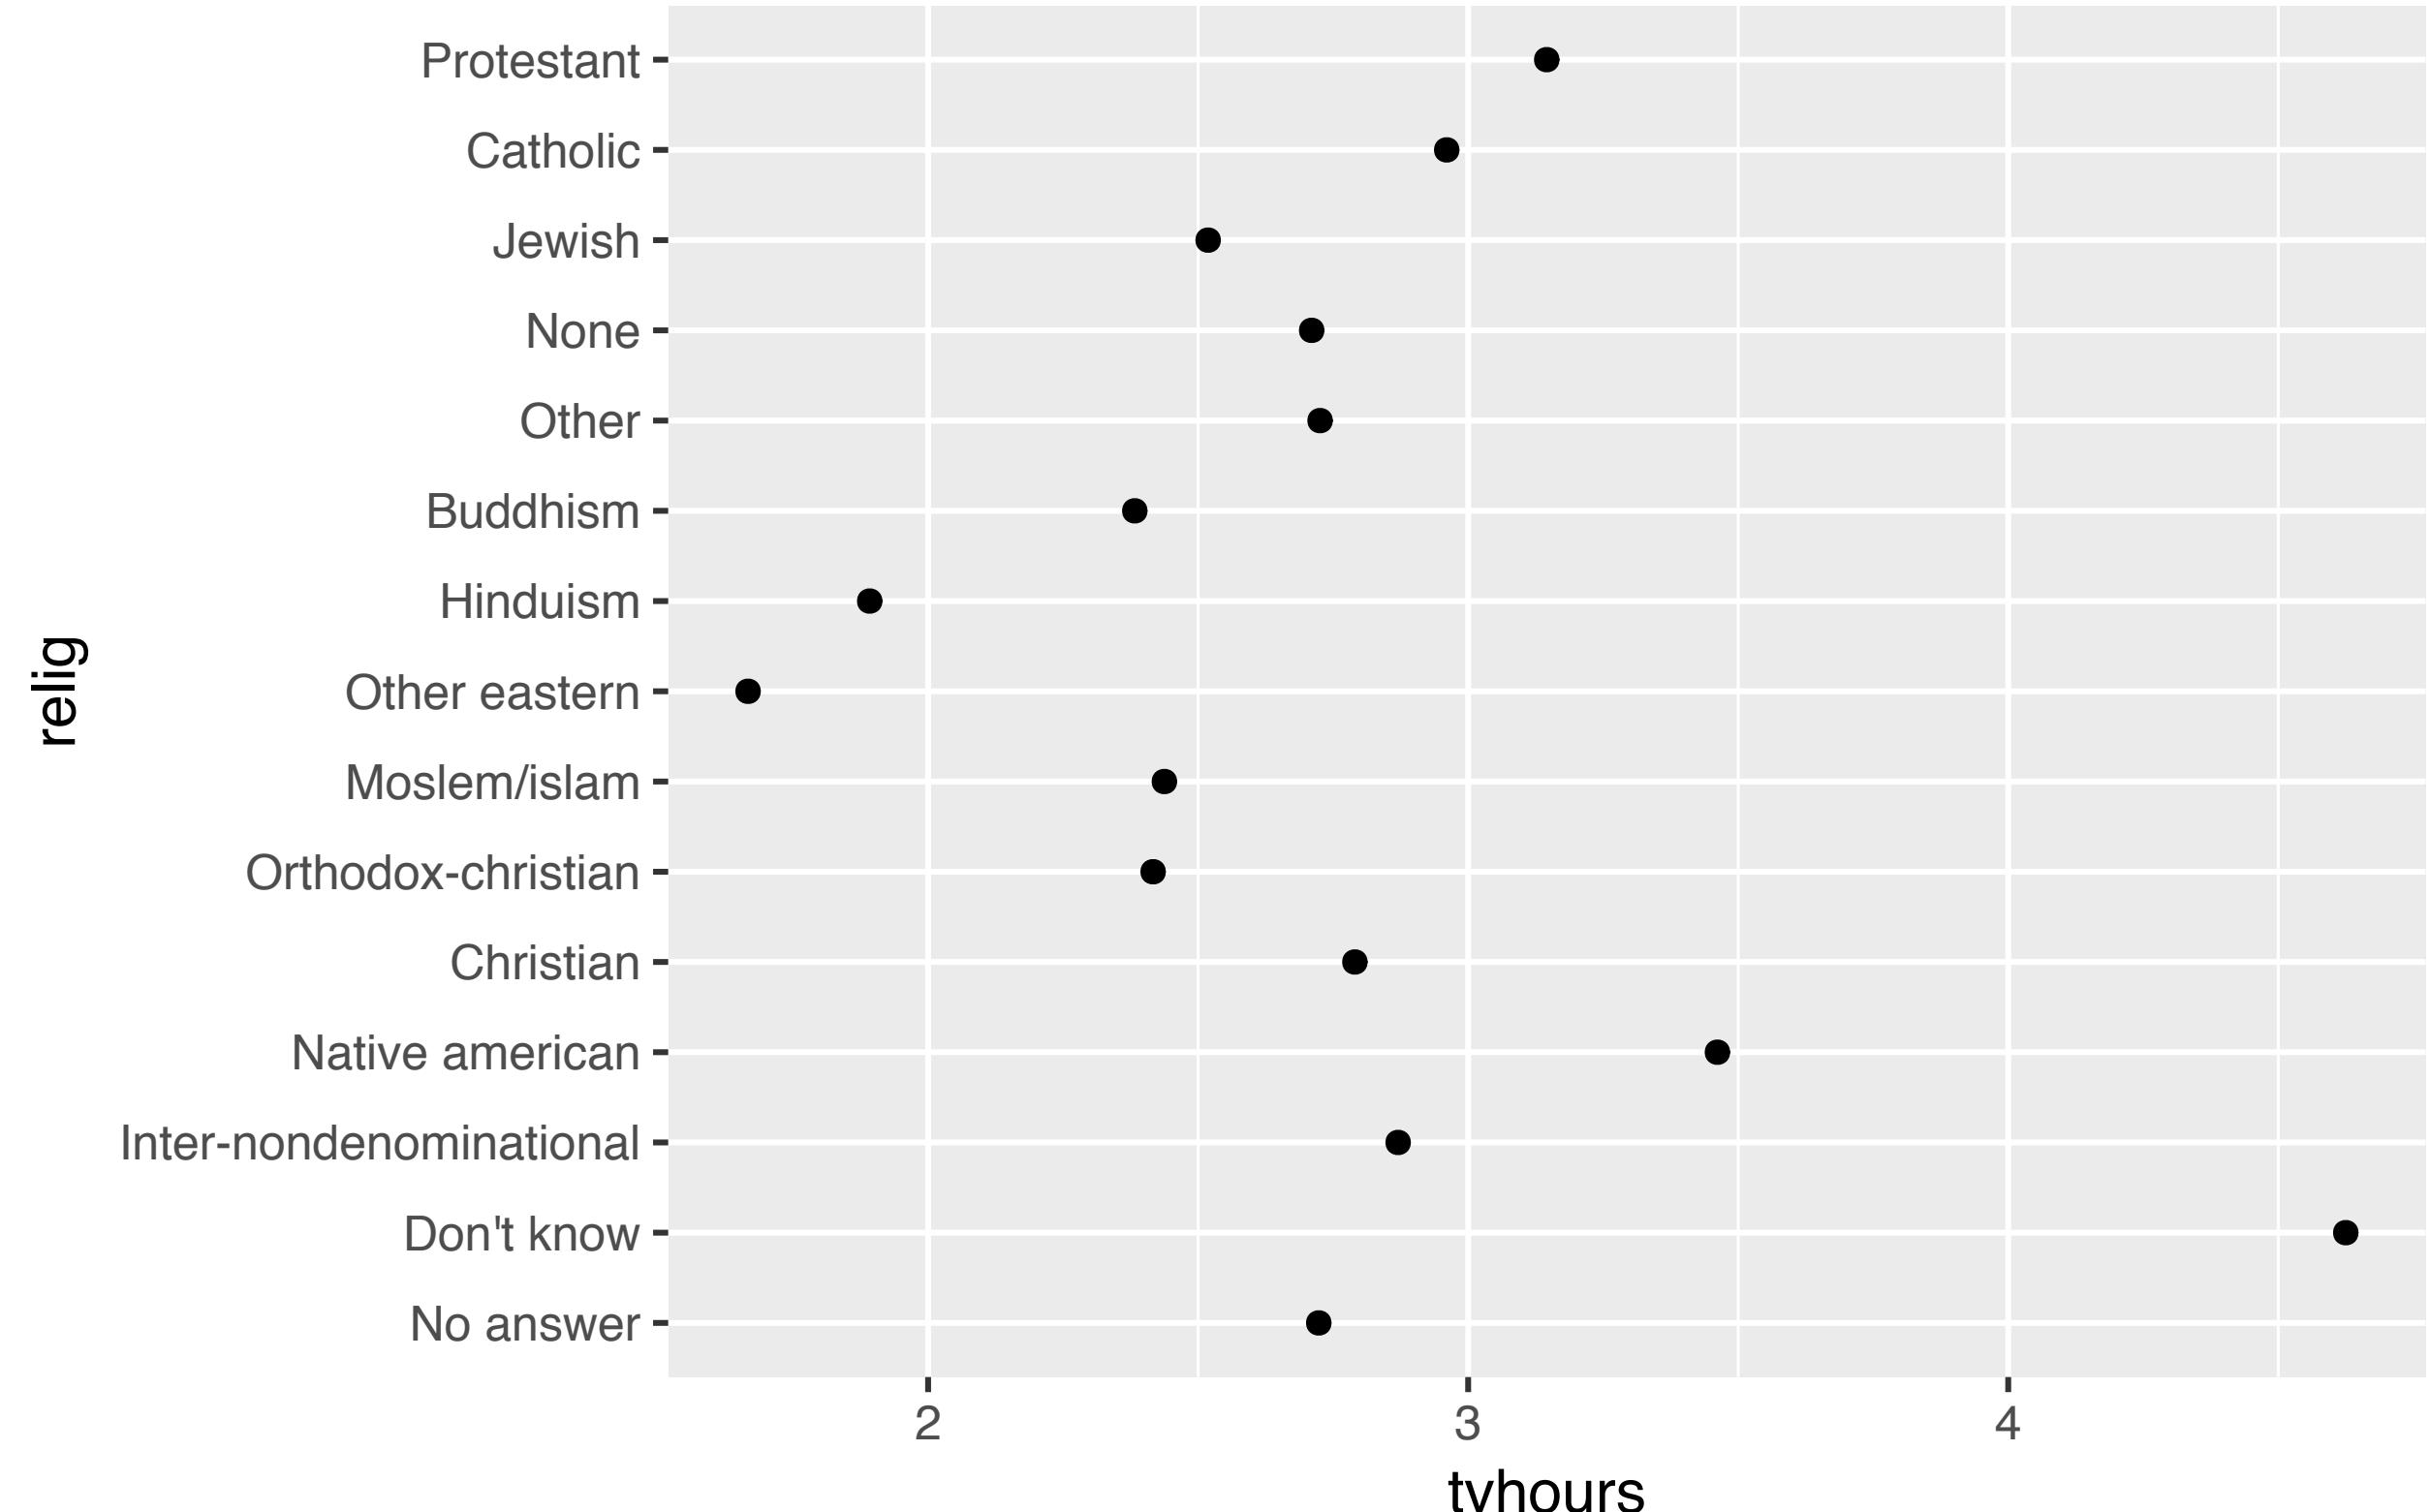


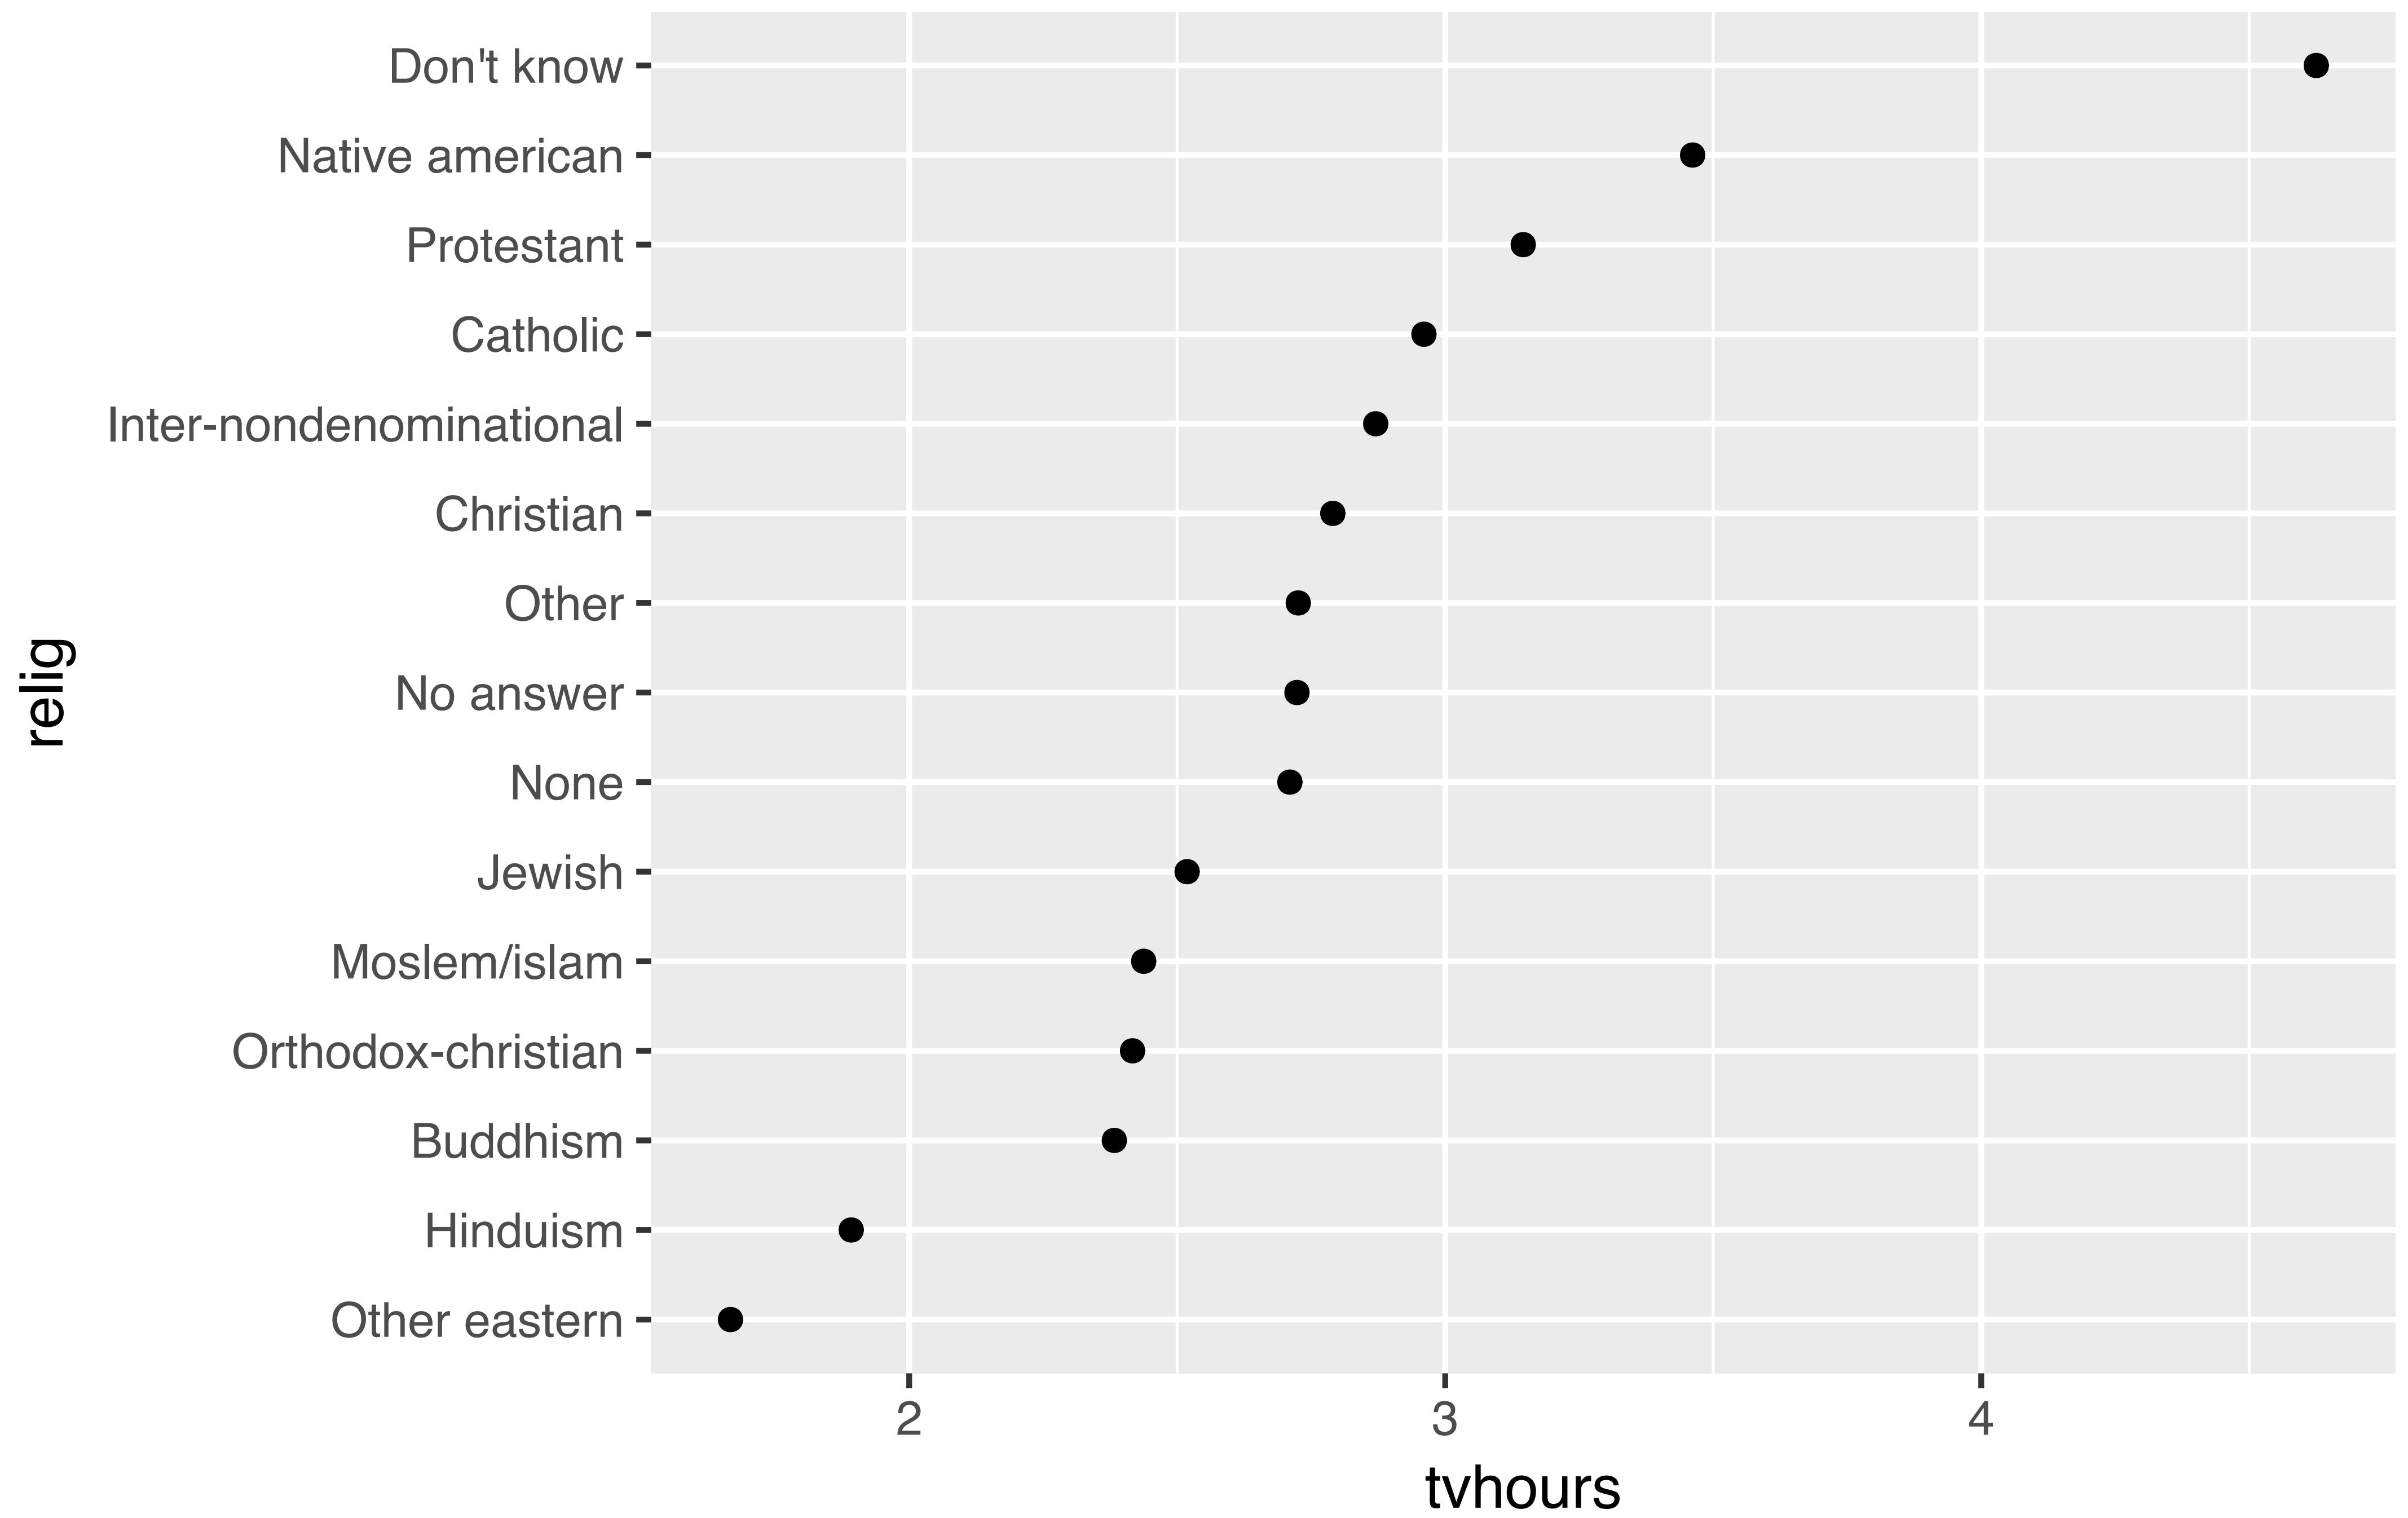
Which religions watch the least TV?

```
gss_cat %>%  
  drop_na(tvhours) %>%  
  group_by(relig) %>%  
  summarise(tvhours = mean(tvhours)) %>%  
  ggplot(aes(tvhours, relig)) +  
    geom_point()
```



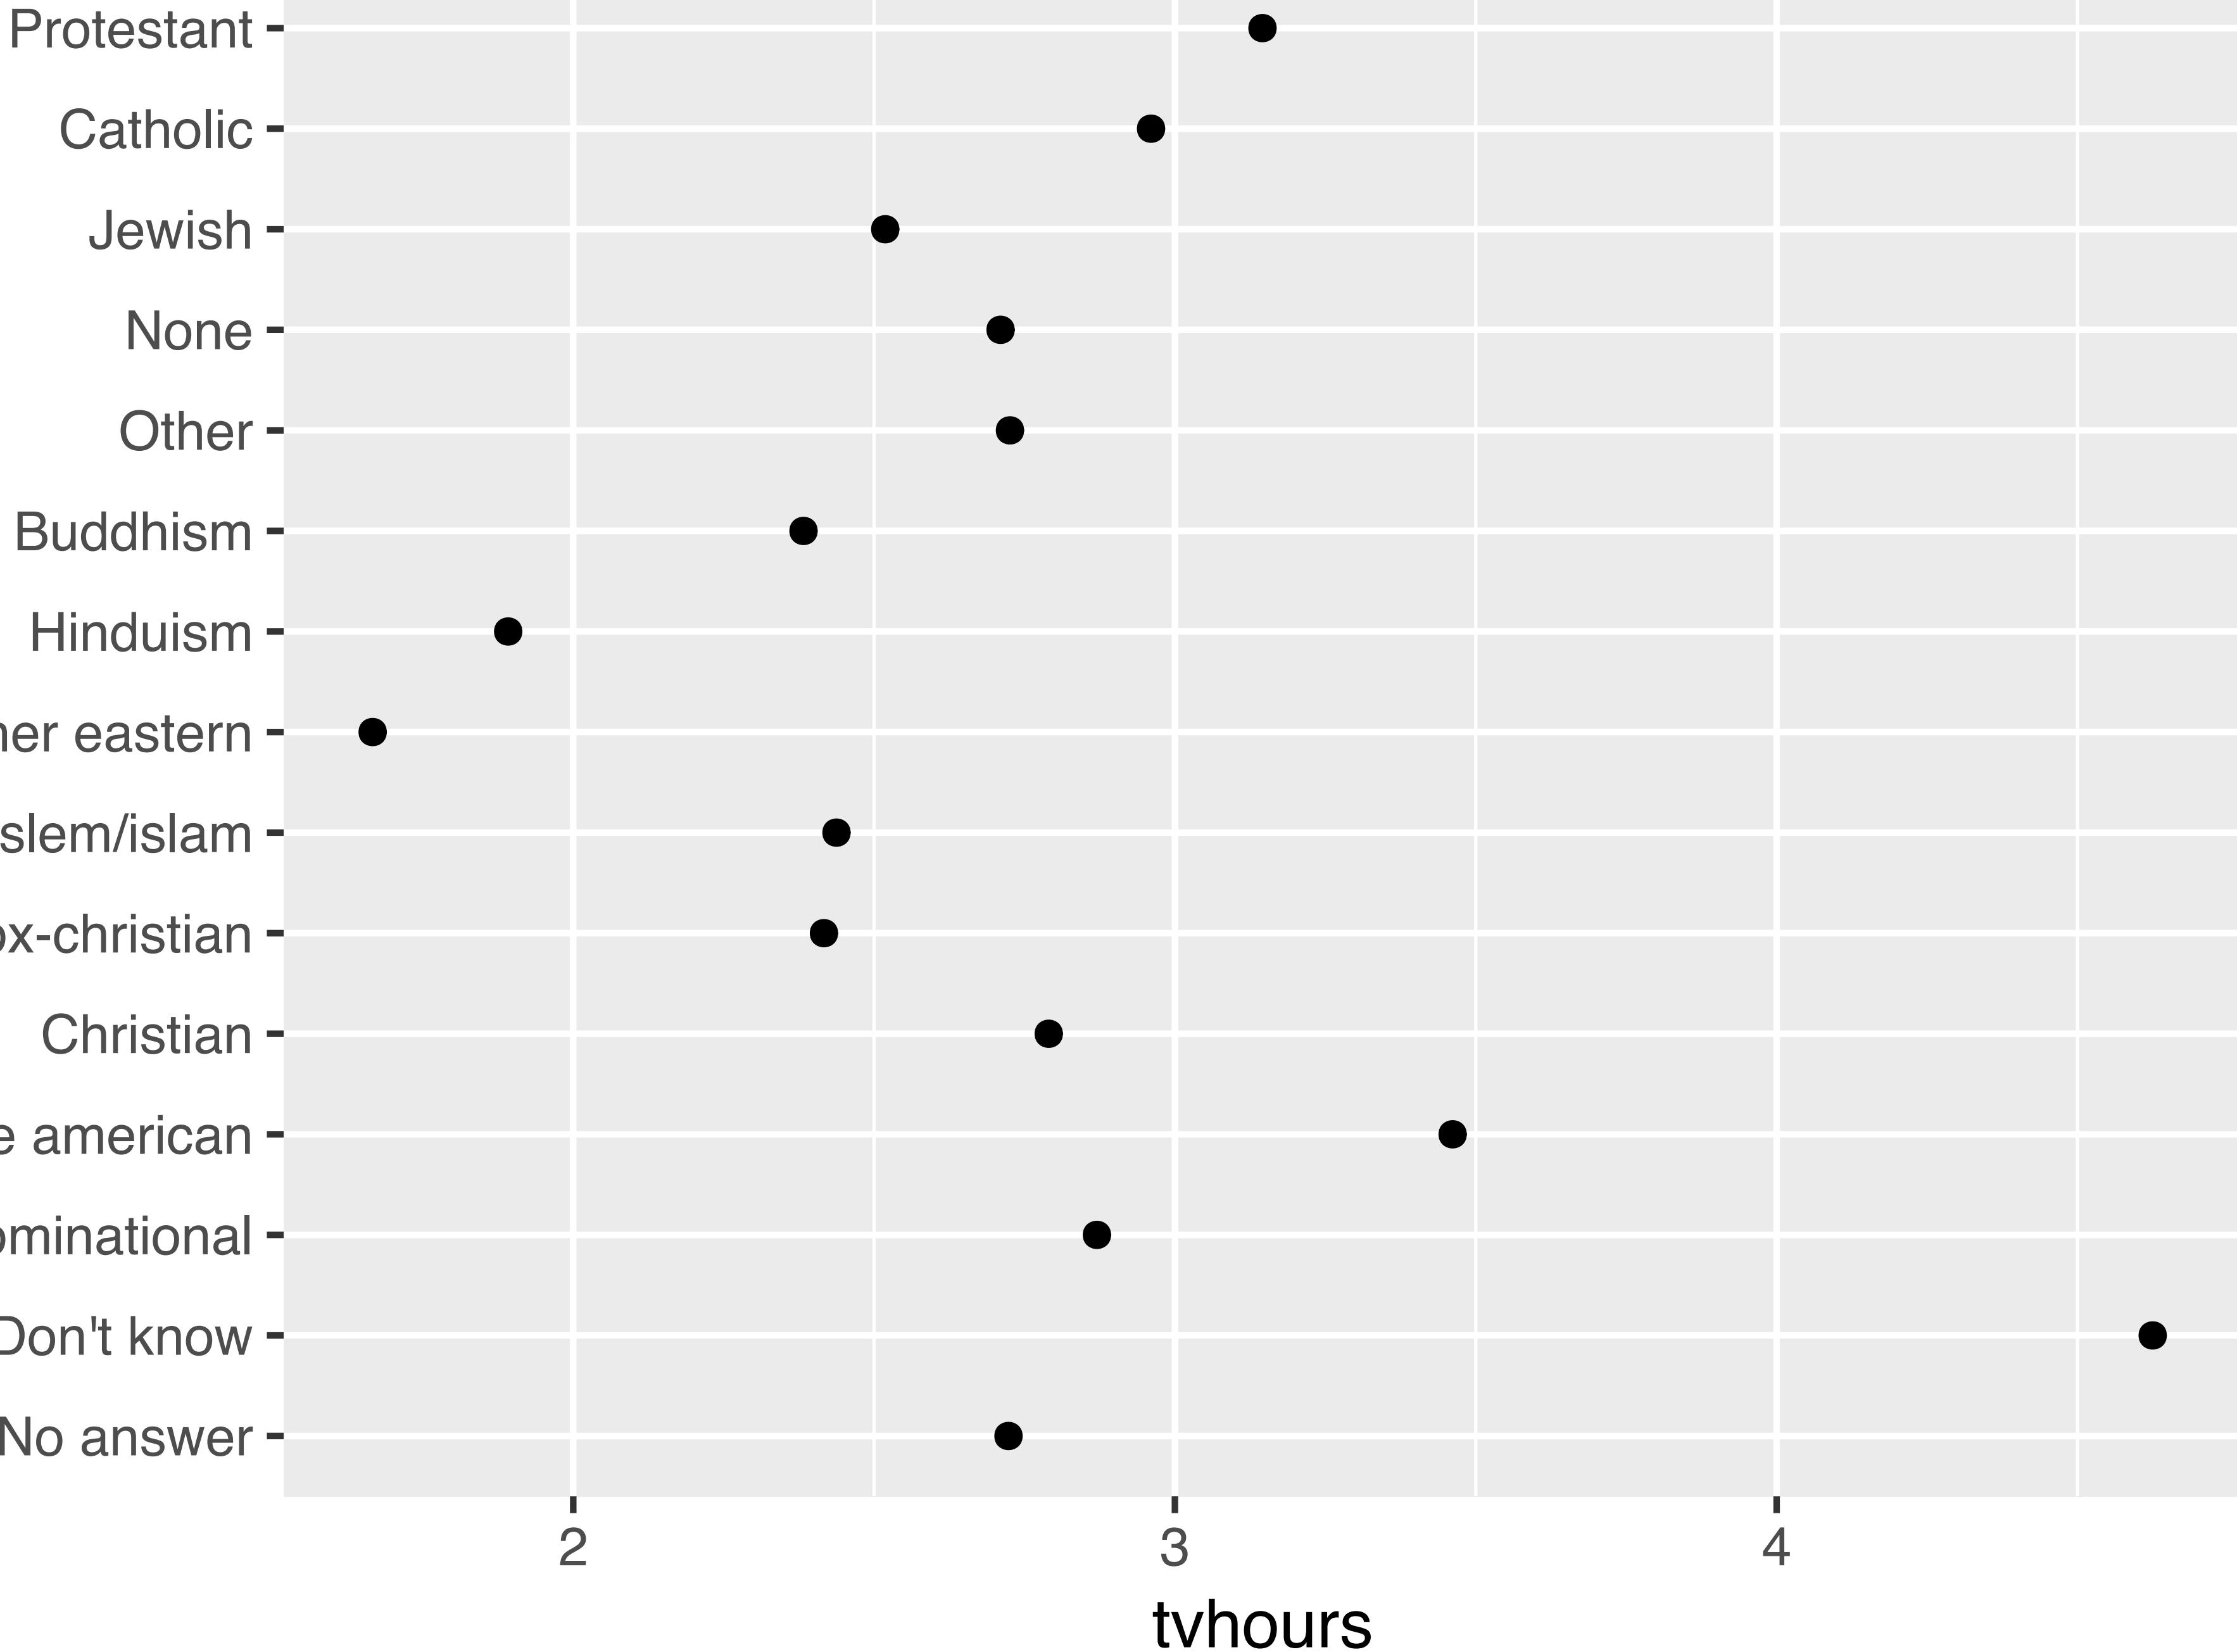
Which do you prefer?





Why is the Y axis in this order?

relig



levels()

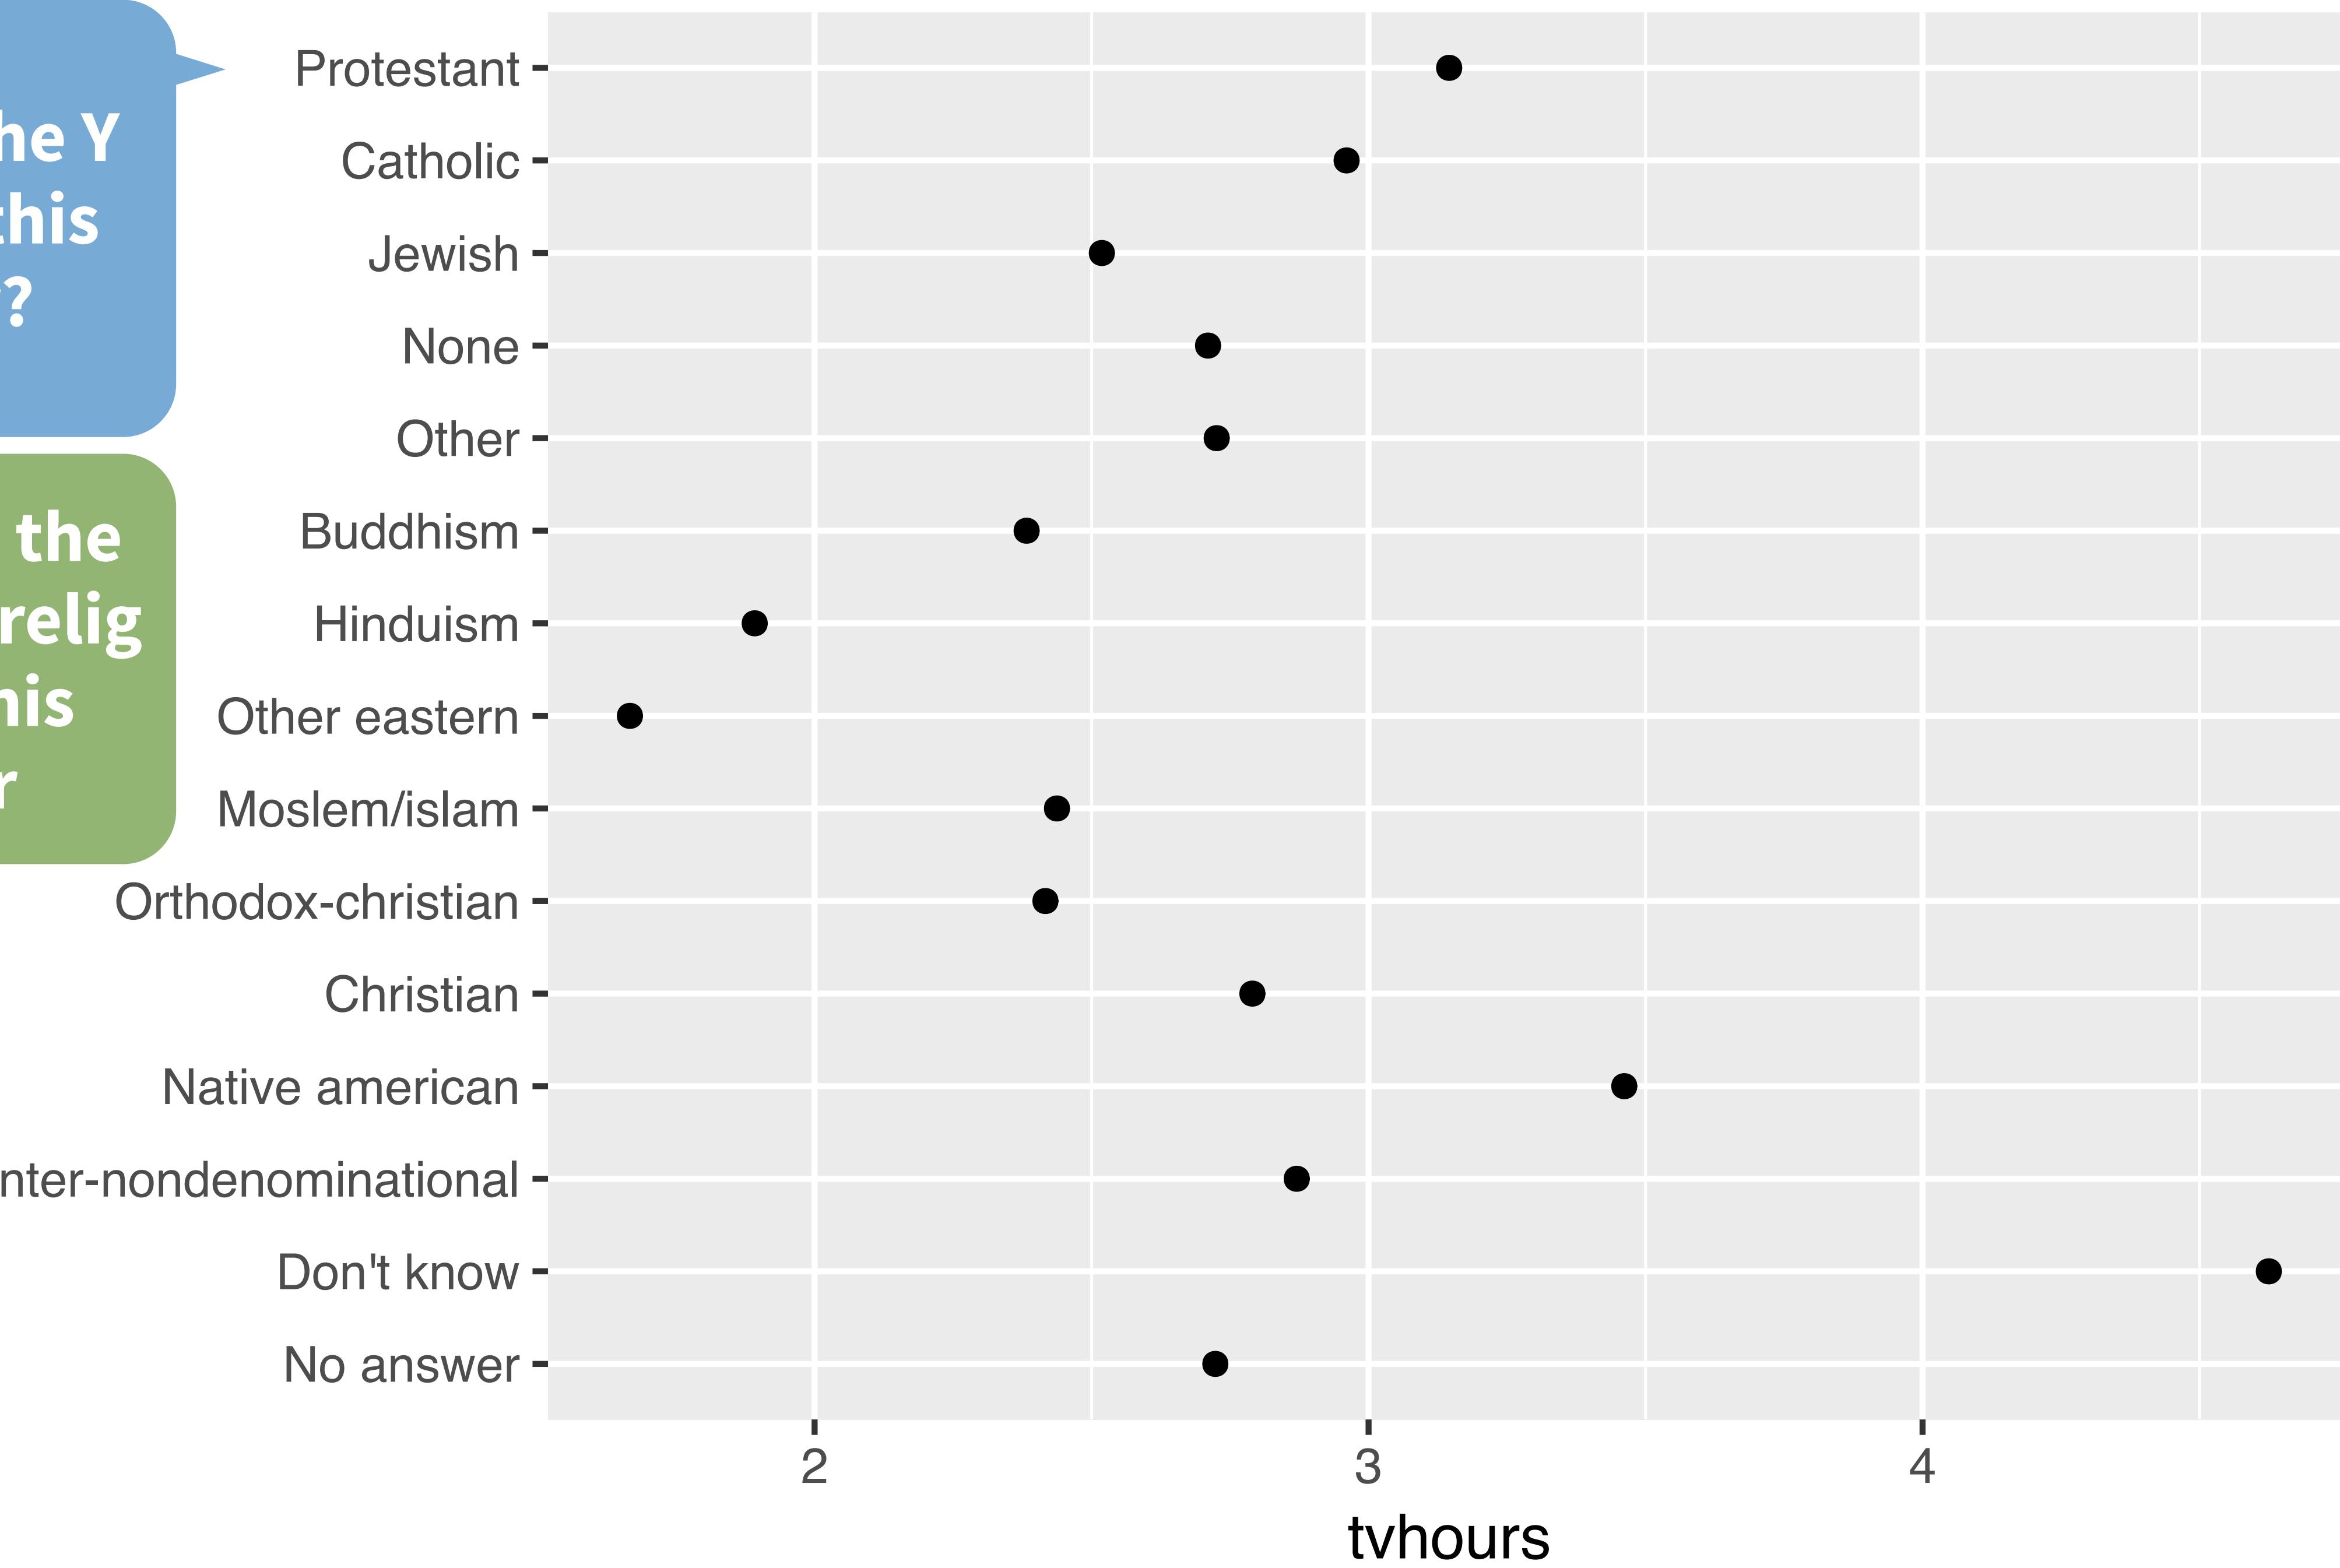
Use **levels()** to access a factor's levels

```
levels(gss_cat$relig)
## [1] "No answer"                      "Don't know"
## [3] "Inter-nondenominational" "Native american"
## [5] "Christian"                      "Orthodox-christian"
## [7] "Moslem/islam"                   "Other eastern"
## [9] "Hinduism"                       "Buddhism"
## [11] "Other"                           "None"
## [13] "Jewish"                          "Catholic"
## [15] "Protestant"                     "Not applicable"
```



Why is the Y axis in this order?

Because the levels of relig have this order



Most useful skills

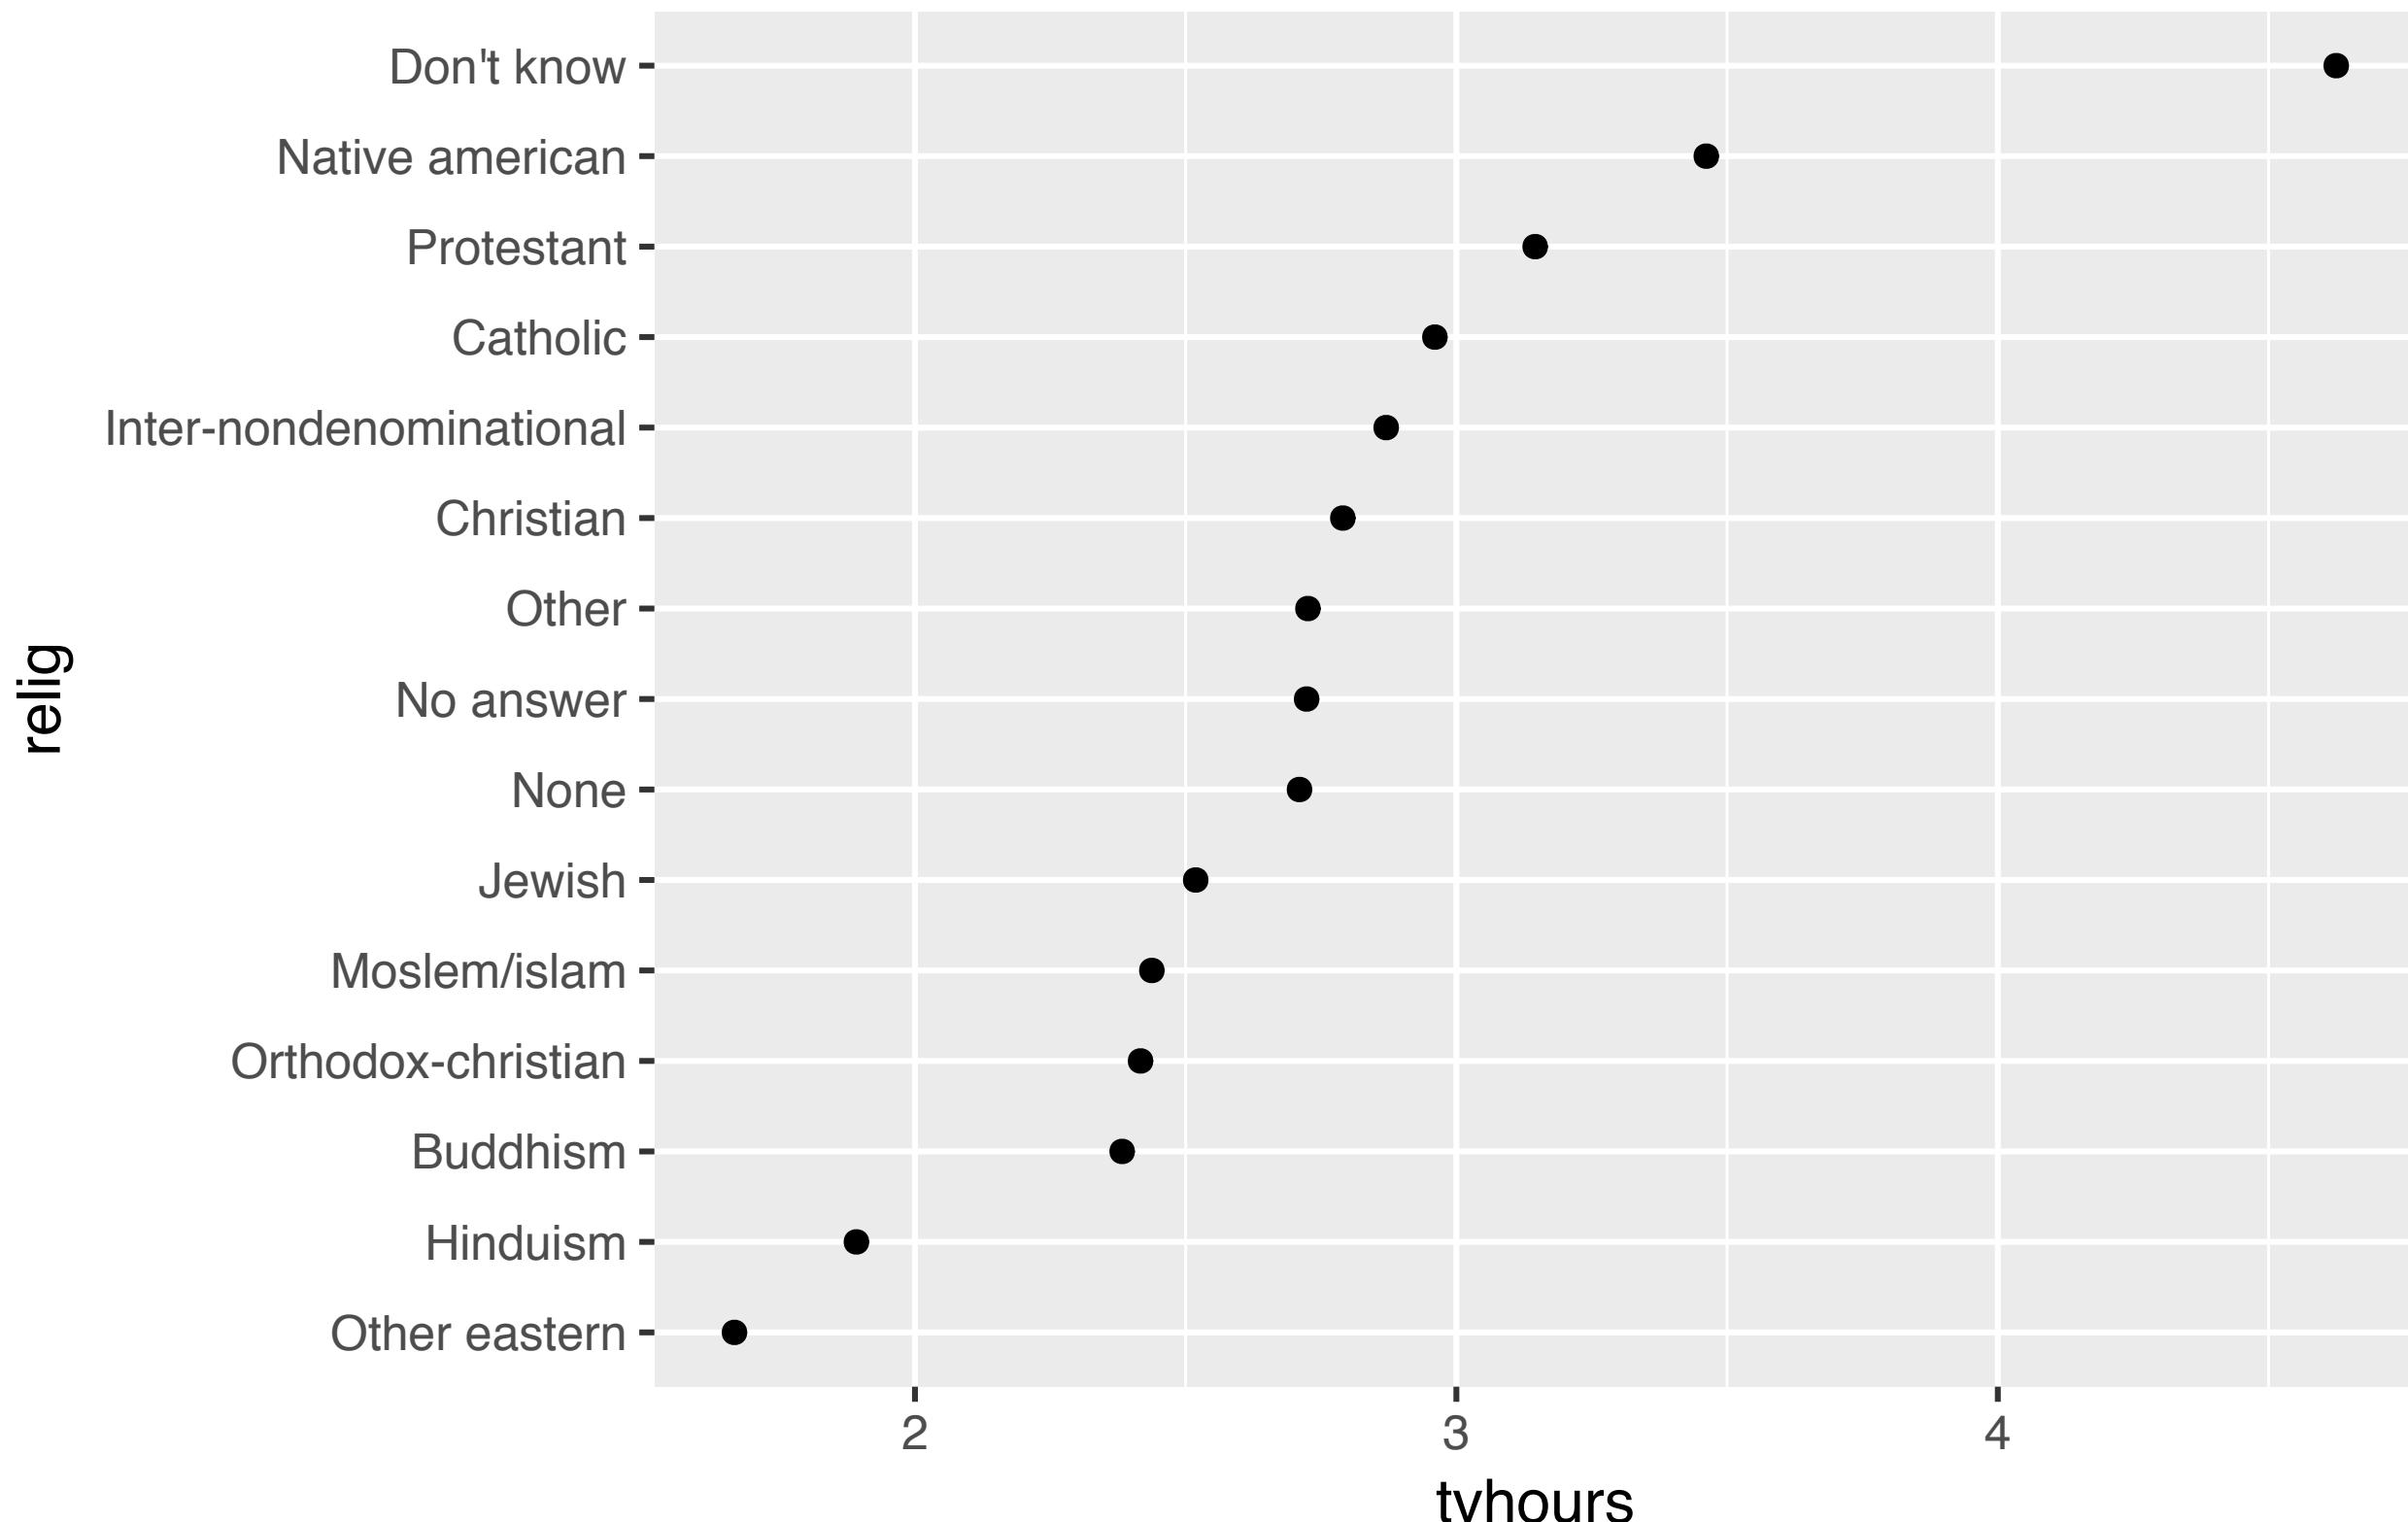
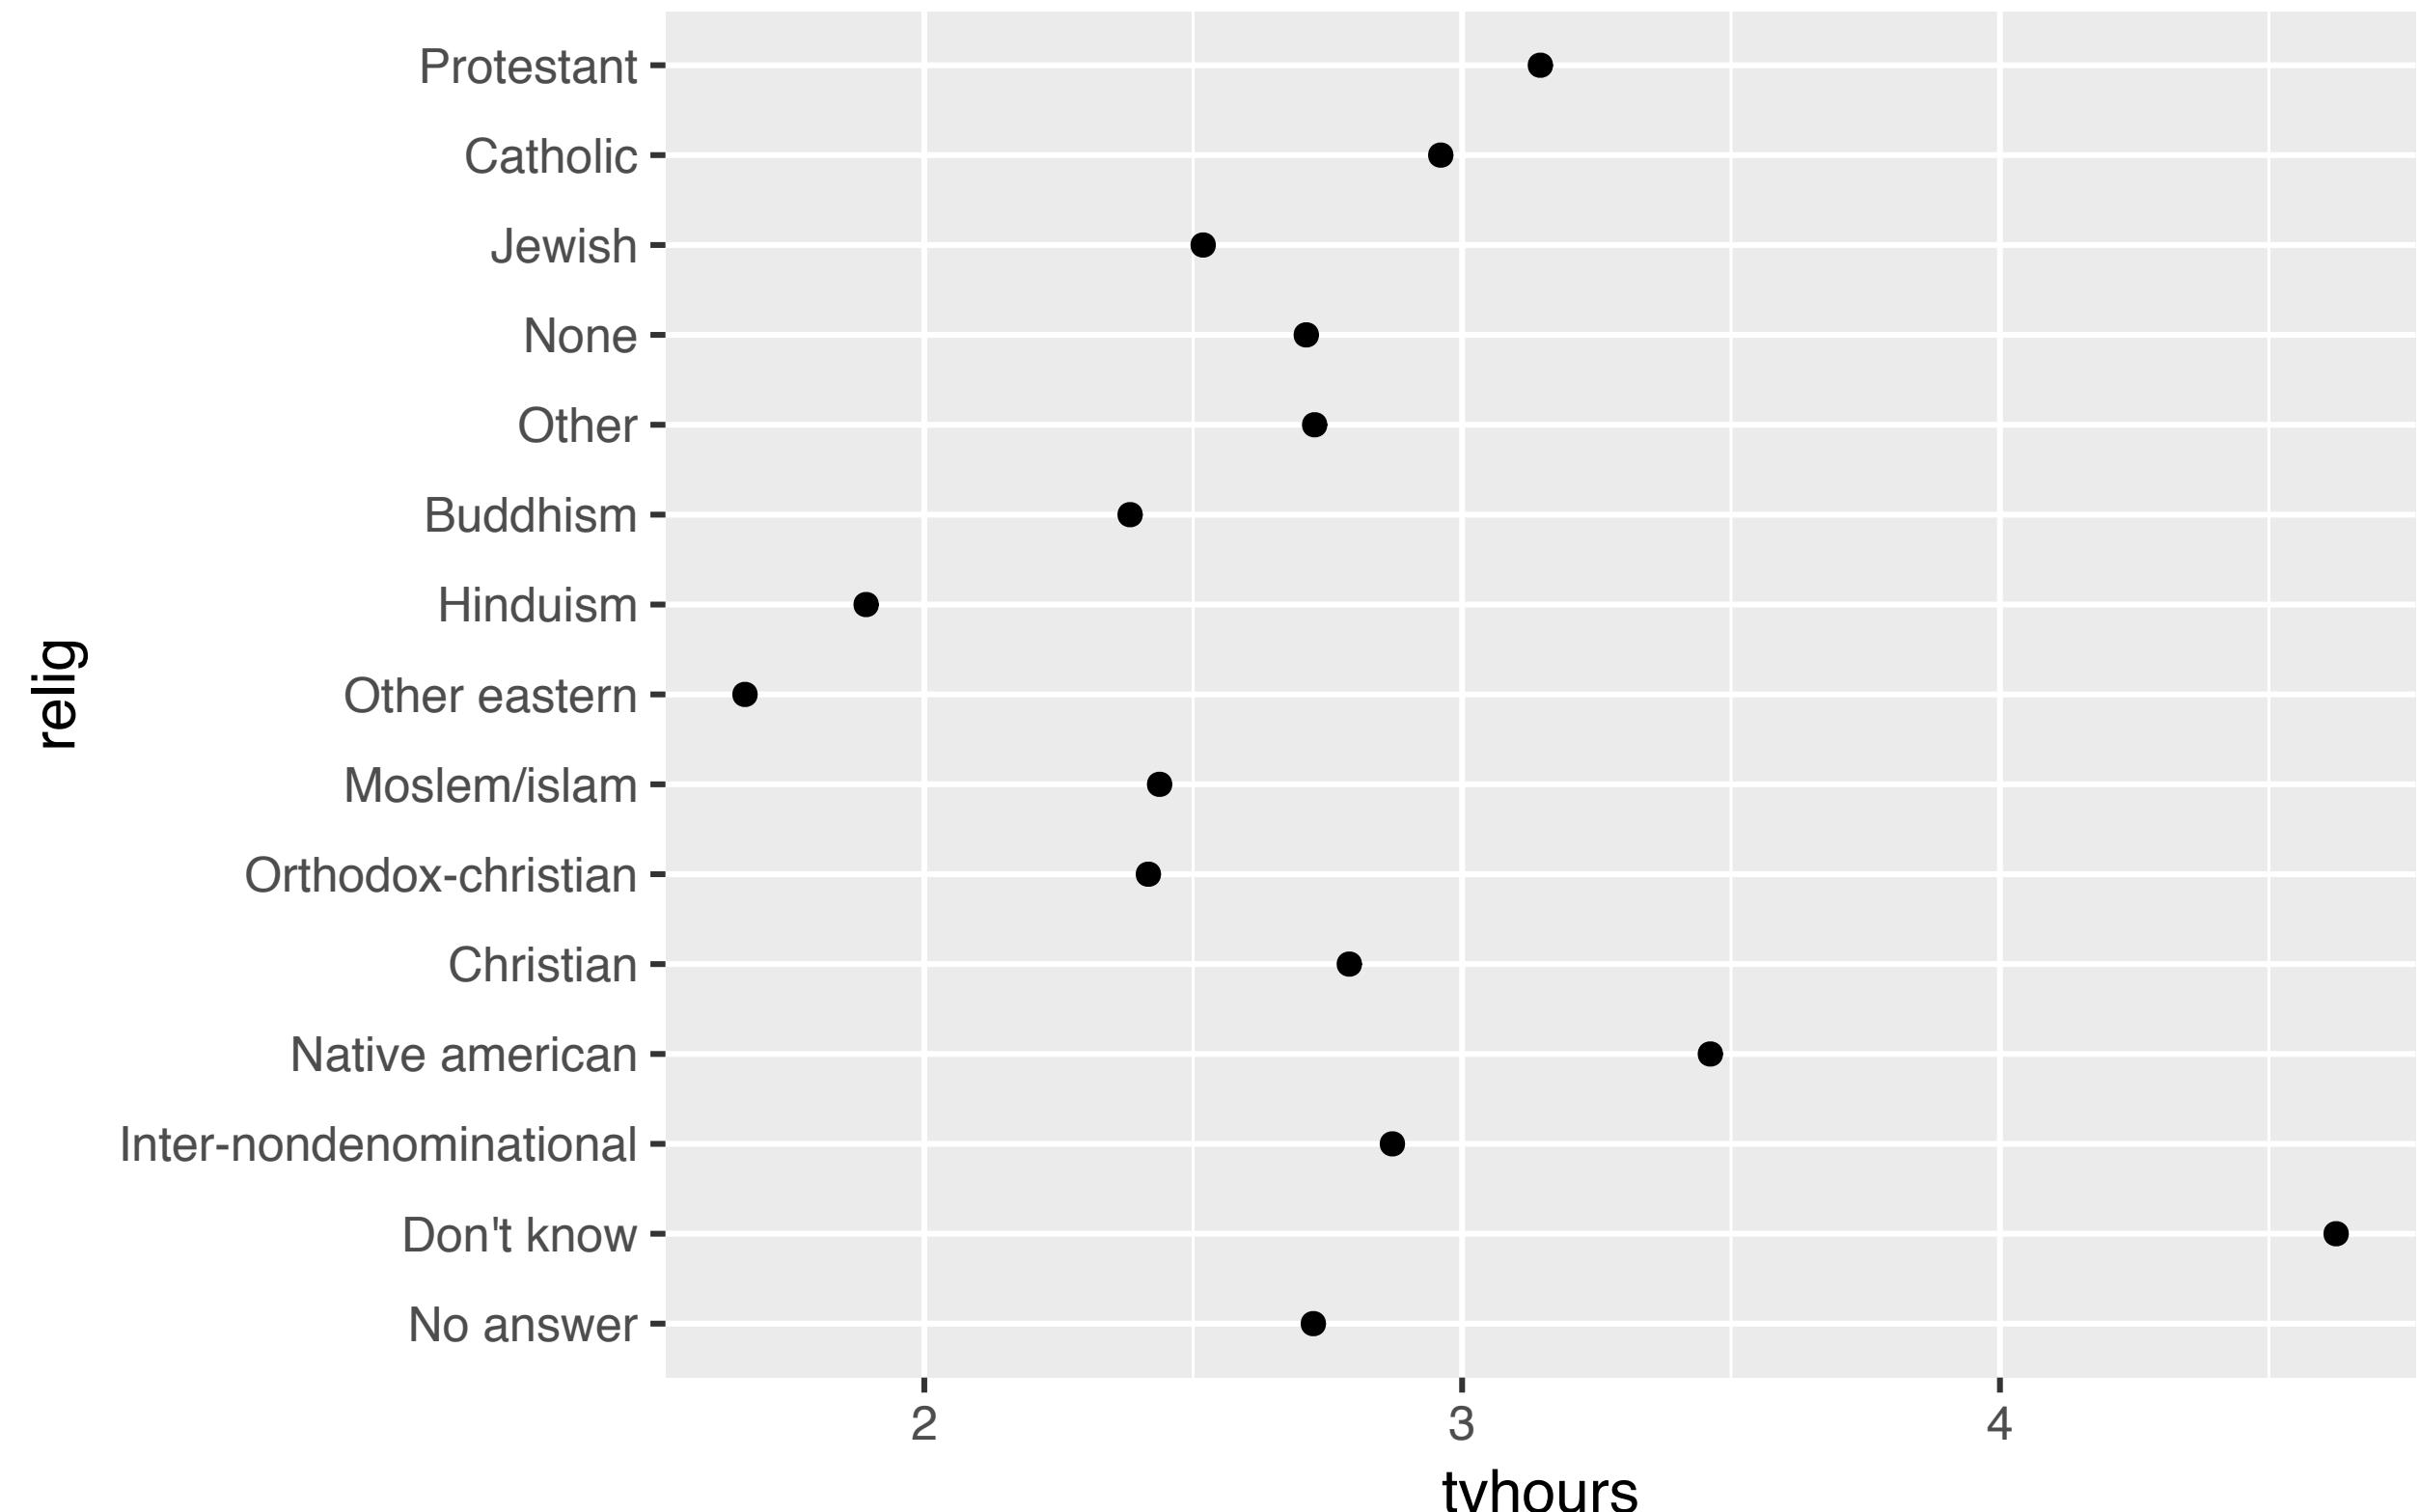
1. Reorder the levels
2. Recode the levels
3. Collapse levels



Reordering levels

A large, semi-transparent watermark of the R logo is positioned in the bottom right corner of the slide. The logo consists of a circular arrow pointing clockwise, with the letters 'R' and 'S' stacked vertically in the center.

Reorder relig by tvhours



fct_reorder()

Reorders the levels of a factor based on the result of `fun(x)` applied to each group of cases (grouped by level).

```
fct_reorder(f, x, fun = median, ..., .desc = FALSE)
```

factor to
reorder

variable to
reorder by
(in conjunction
with fun)

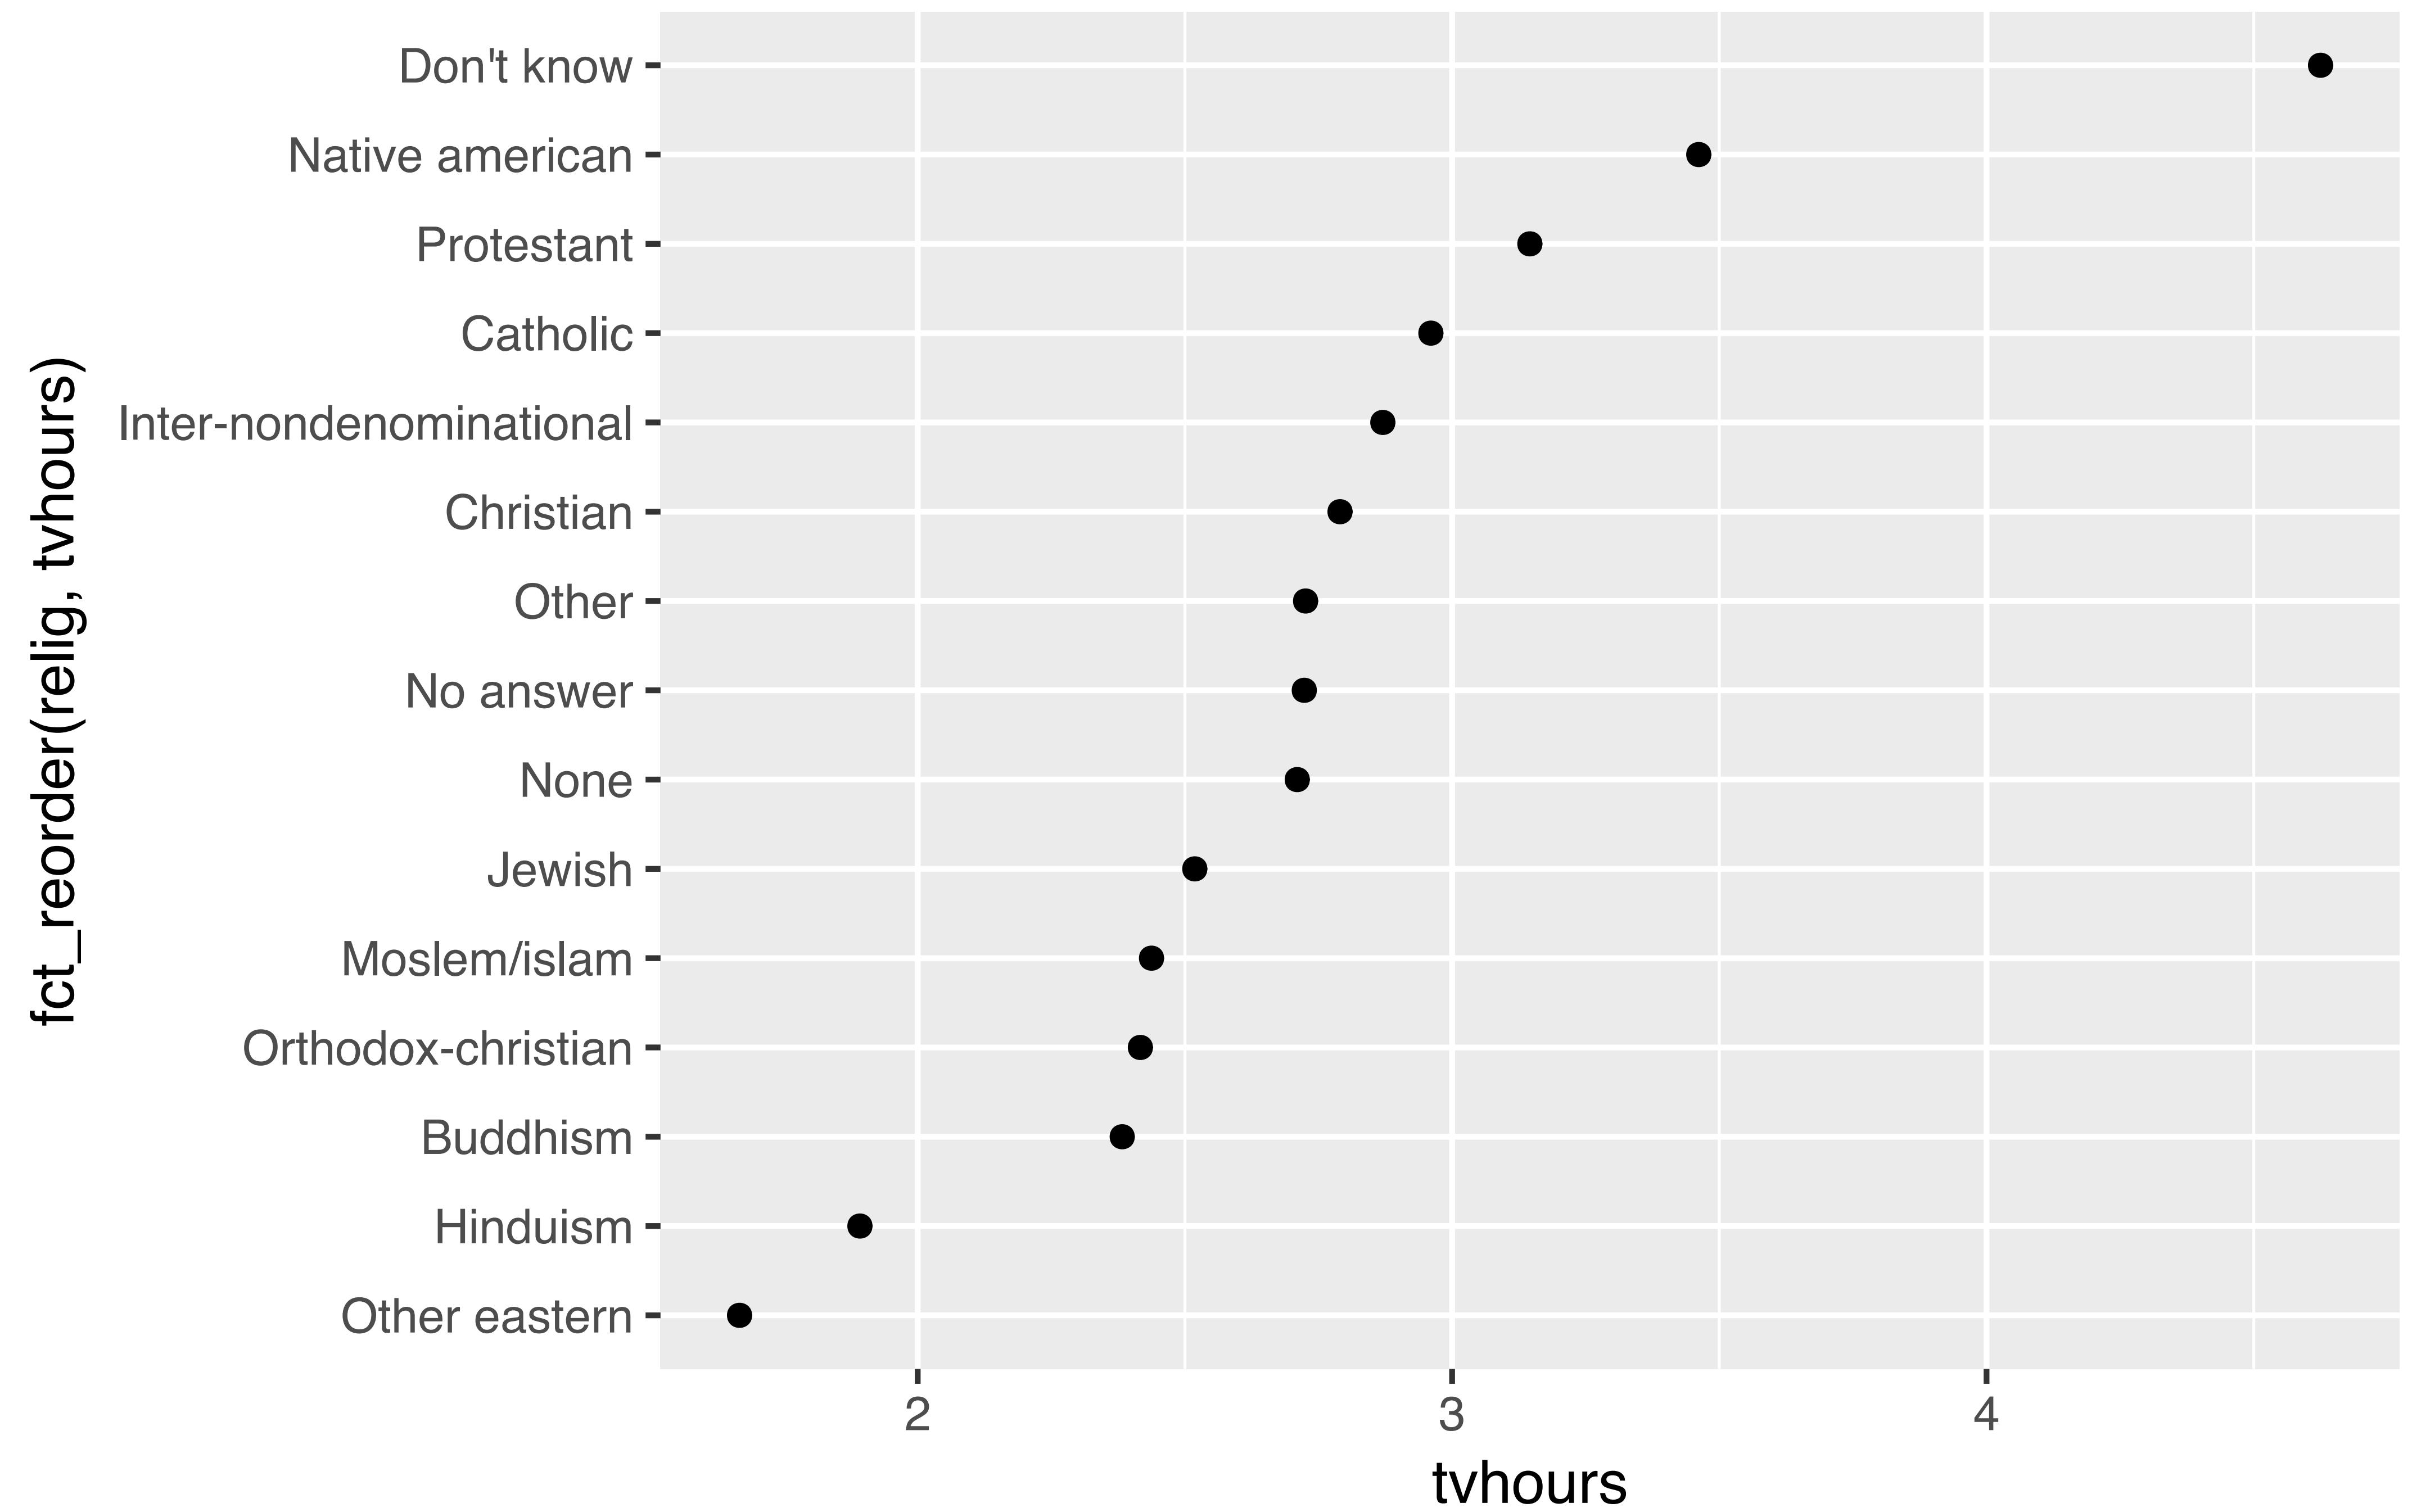
function to
reorder by
(in conjunction
with x)

put in descending
order?



```
gss_cat %>%  
  drop_na(tvhours) %>%  
  group_by(relig) %>%  
  summarise(tvhours = mean(tvhours)) %>%  
  ggplot(aes(tvhours, fct_reorder(relig, tvhours))) +  
  geom_point()
```





Most useful skills

1. Reorder the levels
2. Recode the levels
3. Collapse levels



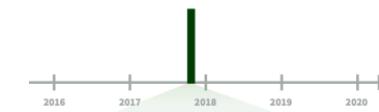
Date times

R

Dates and times with lubridate :: CHEAT SHEET



Date-times



2017-11-28 12:00:00

PARSE DATE-TIMES (Convert strings or numbers to date-times)

1. Identify the order of the year (**y**), month (**m**), day (**d**), hour (**h**), minute (**m**) and second (**s**) elements in your data.
2. Use the function below whose name replicates the order. Each accepts a wide variety of input formats.

2017-11-28T14:02:00	ymd_hms(), ymd_hm(), ymd_h(). ymd_hms("2017-11-28T14:02:00")
2017-22-12 10:00:00	ydm_hms(), ydm_hm(), ydm_h(). ydm_hms("2017-22-12 10:00:00")
11/28/2017 1:02:03	mdy_hms(), mdy_hm(), mdy_h(). mdy_hms("11/28/2017 1:02:03")
1 Jan 2017 23:59:59	dmy_hms(), dmy_hm(), dmy_h(). dmy_hms("1 Jan 2017 23:59:59")
20170131	ymd(), ydm(). ymd(20170131)
July 4th, 2000	mdy(), myd(). mdy("July 4th, 2000")
4th of July '99	dmy(), dyd(). dmy("4th of July '99")
2001: Q3	yq() Q for quarter. yq("2001: Q3")
2.01	hms::hms() Also lubridate::hms(), lubr::hms() and lubridate::timespan()

2017-11-28 12:00:00

A **date-time** is a point on the timeline, stored as the number of seconds since 1970-01-01 00:00:00 UTC

```
dt <- as_datetime(1511870400)  
## "2017-11-28 12:00:00 UTC"
```

2017-11-28

A **date** is a day stored as the number of days since 1970-01-01

```
d <- as_date(17498)  
## "2017-11-28"
```

12:00:00

An **hms** is a **time** stored as the number of seconds since 00:00:00

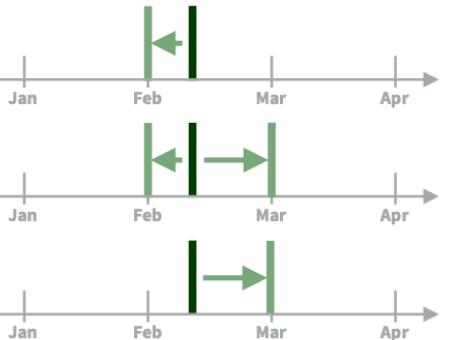
```
t <- hms::as.hms(85)  
## 00:01:25
```

GET AND SET COMPONENTS

Use an accessor function to get a component. Assign into an accessor function to change a component in place.

```
d ## "2017-11-28"  
day(d) ## 28  
day(d) <- 1  
d ## "2017-11-01"
```

Round Date-times



floor_date(x, unit = "second")
Round down to nearest unit.
floor_date(dt, unit = "month")

round_date(x, unit = "second")
Round to nearest unit.
round_date(dt, unit = "month")

ceiling_date(x, unit = "second", change_on_boundary = NULL)
Round up to nearest unit.
ceiling_date(dt, unit = "month")

rollback(dates, roll_to_first = FALSE, preserve_hms = TRUE)
Roll back to last day of previous month. **rollback(dt)**

Stamp Date-times

stamp() Derive a template from an example string and return a new function that will apply the template to date-times. Also **stamp_date()** and **stamp_time()**.

1. Derive a template, create a function
`sf <- stamp("Created Sunday, Jan 17, 1999 3:34")`
2. Apply the template to dates
`sf(ymd("2010-04-05"))
[1] "Created Monday, Apr 05, 2010 00:00"`

Tip: use a date with day > 12

Time Zones

R recognizes ~600 time zones. Each encodes the time zone, Daylight Savings Time, and historical calendar variations for an area. Preassigned

lubridate Cheat Sheet

Please take out

Quiz

Does every year have 365 days?

Quiz

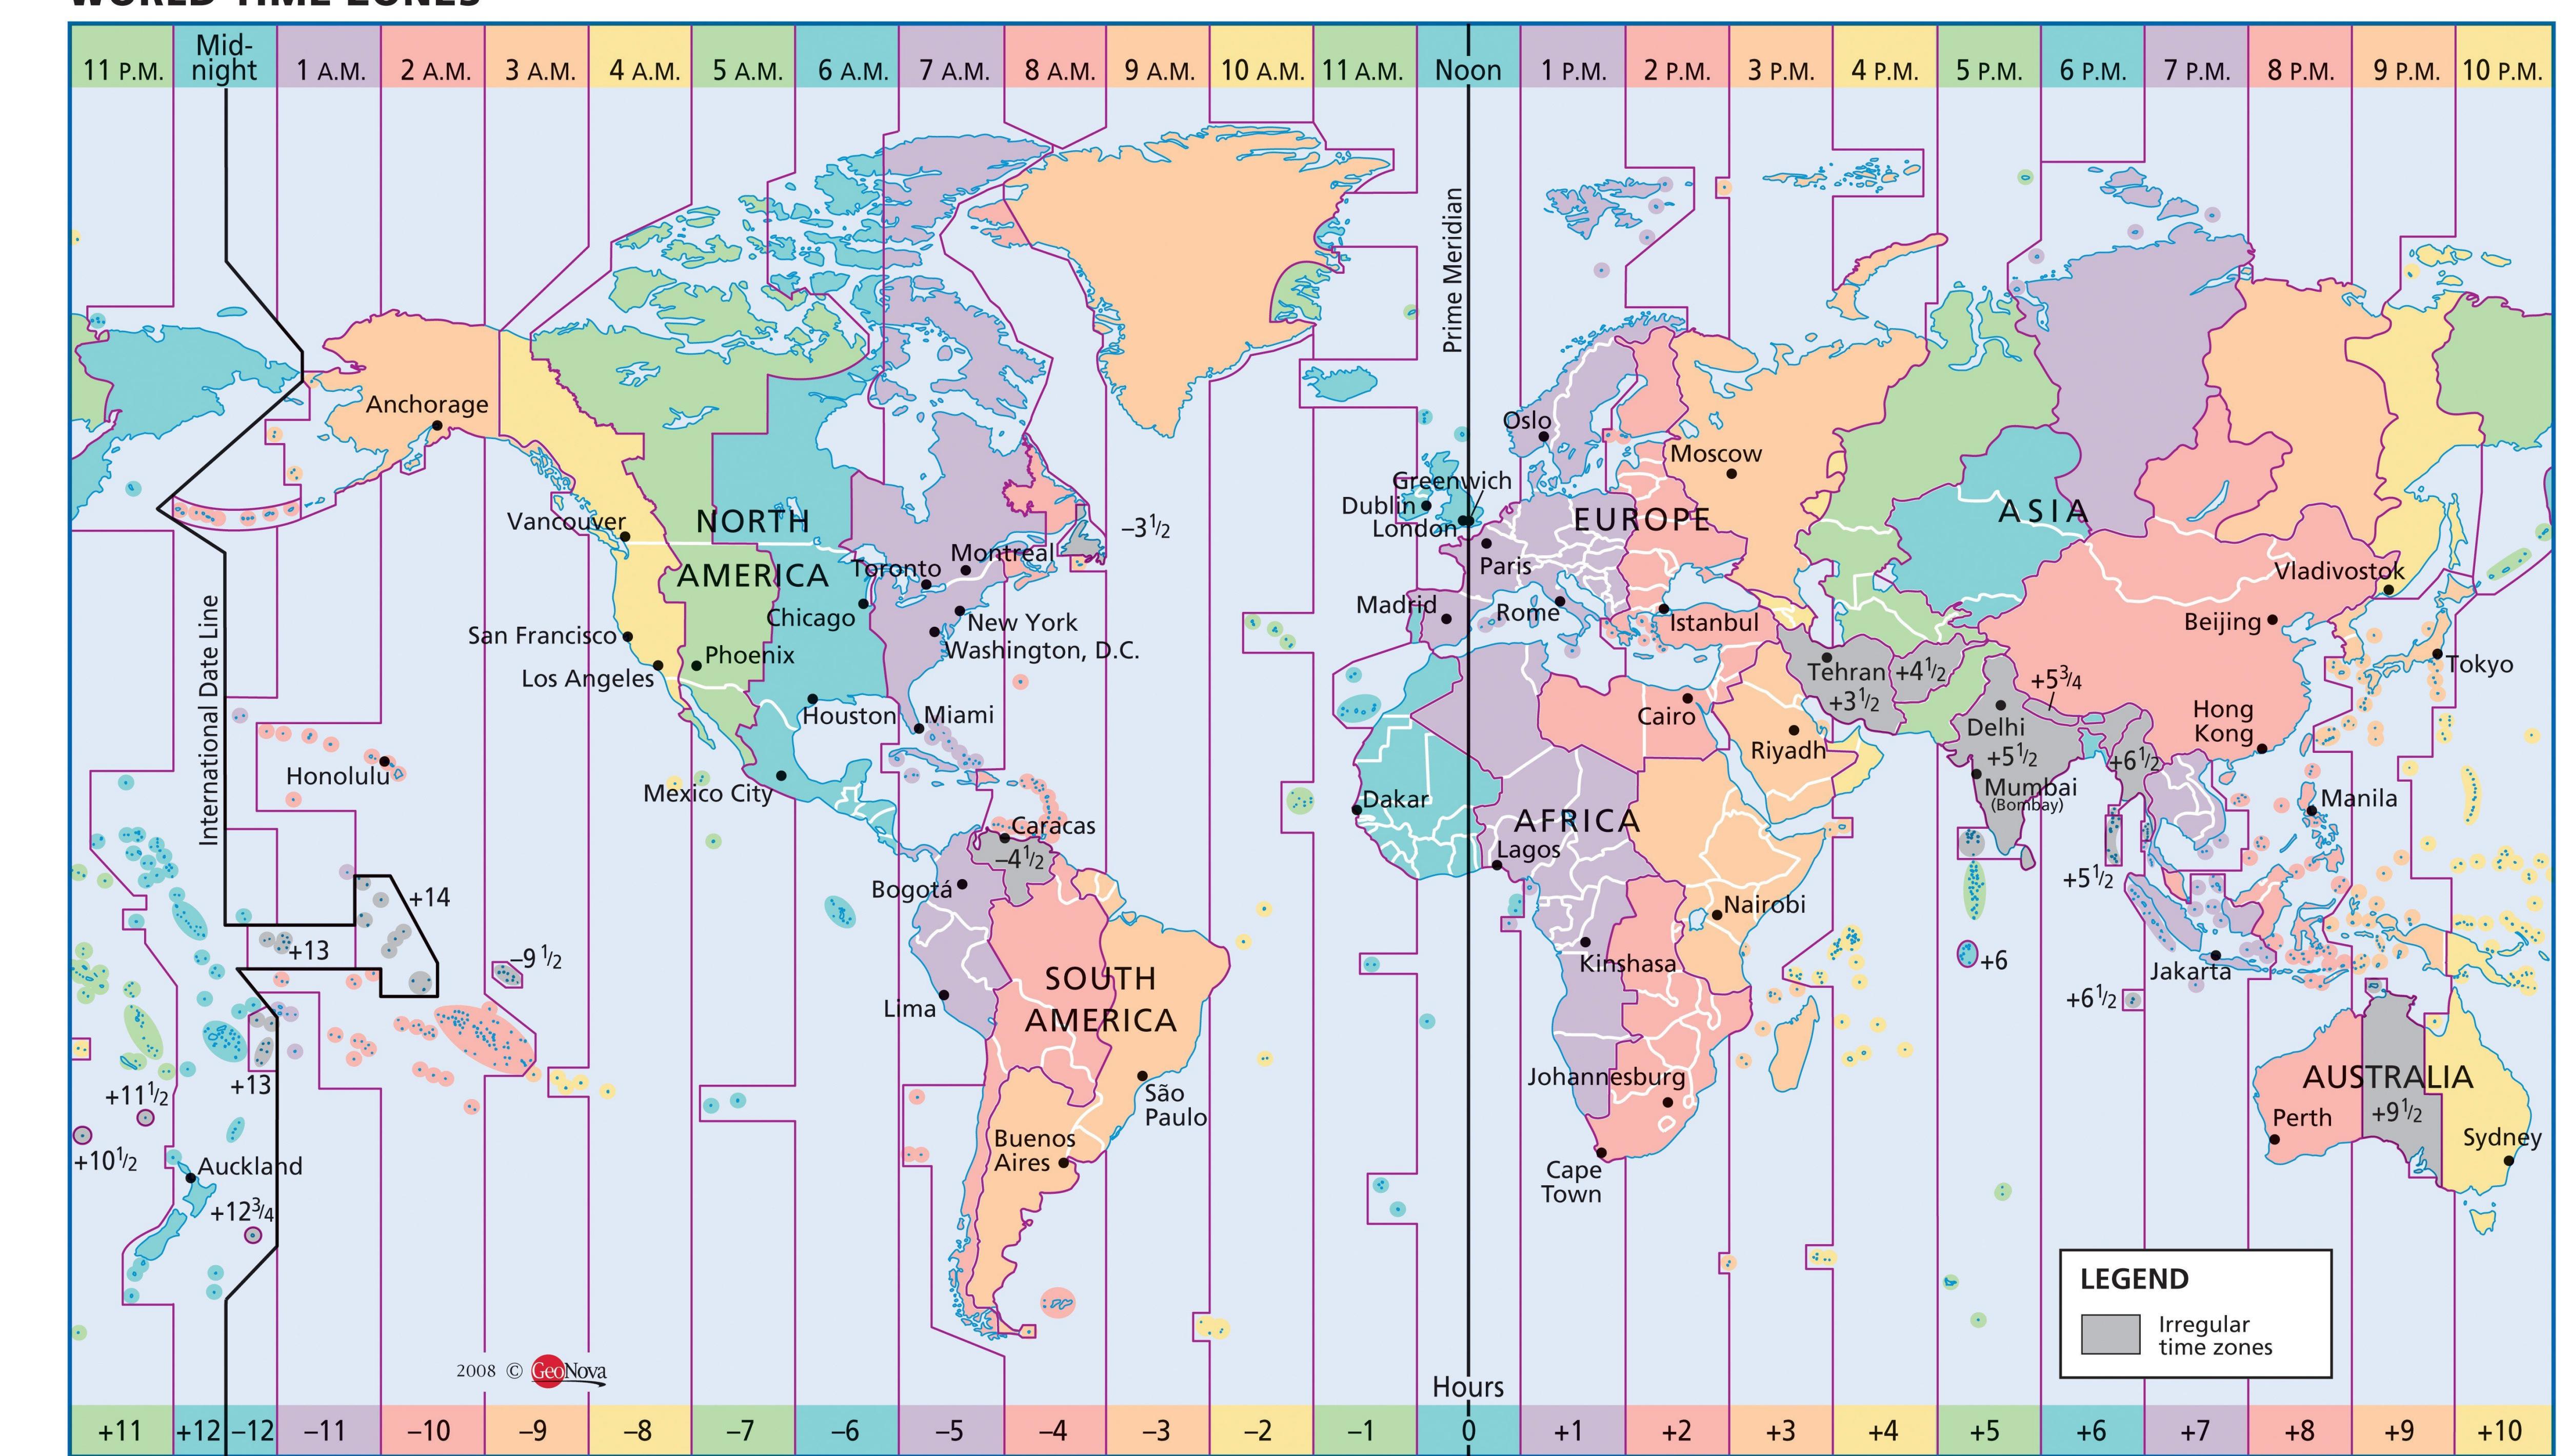
Does every day have 24 hours?

Quiz

Does every minute have 60 seconds?

Quiz

What does a month measure?



Most useful skills

1. Creating dates/times (i.e. *parsing*)
2. Access and change parts of a date



Warm Up

Decide in your group:

- What is the best day of the week to fly?



Creating dates and times



hms



A class for representing just clock times.

```
# install.packages("tidyverse")
library(hms)
```



lubridate



Functions for working with dates and time spans

```
# install.packages("tidyverse")
library(lubridate)
```



ymd() family

To parse strings as dates, use a y, m, d, h, m, s combo

```
ymd("2017/01/11")
mdy("January 11, 2017")
ymd_hms("2017-01-11 01:30:55")
```



Parsing functions

function	parses to
ymd_hms(), ymd_hm(), ymd_h()	
ydm_hms(), ydm_hm(), ydm_h()	POSIXct
dmy_hms(), dmy_hm(), dmy_h()	
mdy_hms(), mdy_hm(), mdy_h()	
ymd(), ydm(), mdy()	
myd(), dmy(), dym(), yq()	Date (POSIXct if tz specified)
hms(), hm(), ms()	Period



Accessing
and changing
components



Accessing components

Extract components by name with a **singular** name

```
date <- ymd("2017-01-11")
year(date)
## 2017
```



Setting components

Use the same function to set components

```
date  
## "2017-01-11"  
year(date) <- 1999  
date  
## "1999-01-11"
```



Accessing date time components

function	extracts	extra arguments
year()	year	
month()	month	label = FALSE, abbr = TRUE
week()	week	
day()	day of month	
wday()	day of week	label = FALSE, abbr = TRUE
qday()	day of quarter	
yday()	day of year	
hour()	hour	
minute()	minute	
second()	second	



Accessing components

```
wday(ymd("2017-01-11"))

## 4

wday(ymd("2017-01-11"), label = TRUE)

## [1] Wed

## 7 Levels: Sun < Mon < Tues < Wed < Thurs < ... < Sat

wday(ymd("2017-01-11"), label = TRUE, abbr = FALSE)

## [1] Wednesday

## 7 Levels: Sunday < Monday < Tuesday < ... < Saturday
```



Your Turn 6

Fill in the blanks to:

Extract the day of the week of each flight (as a full name) from **time_hour**.

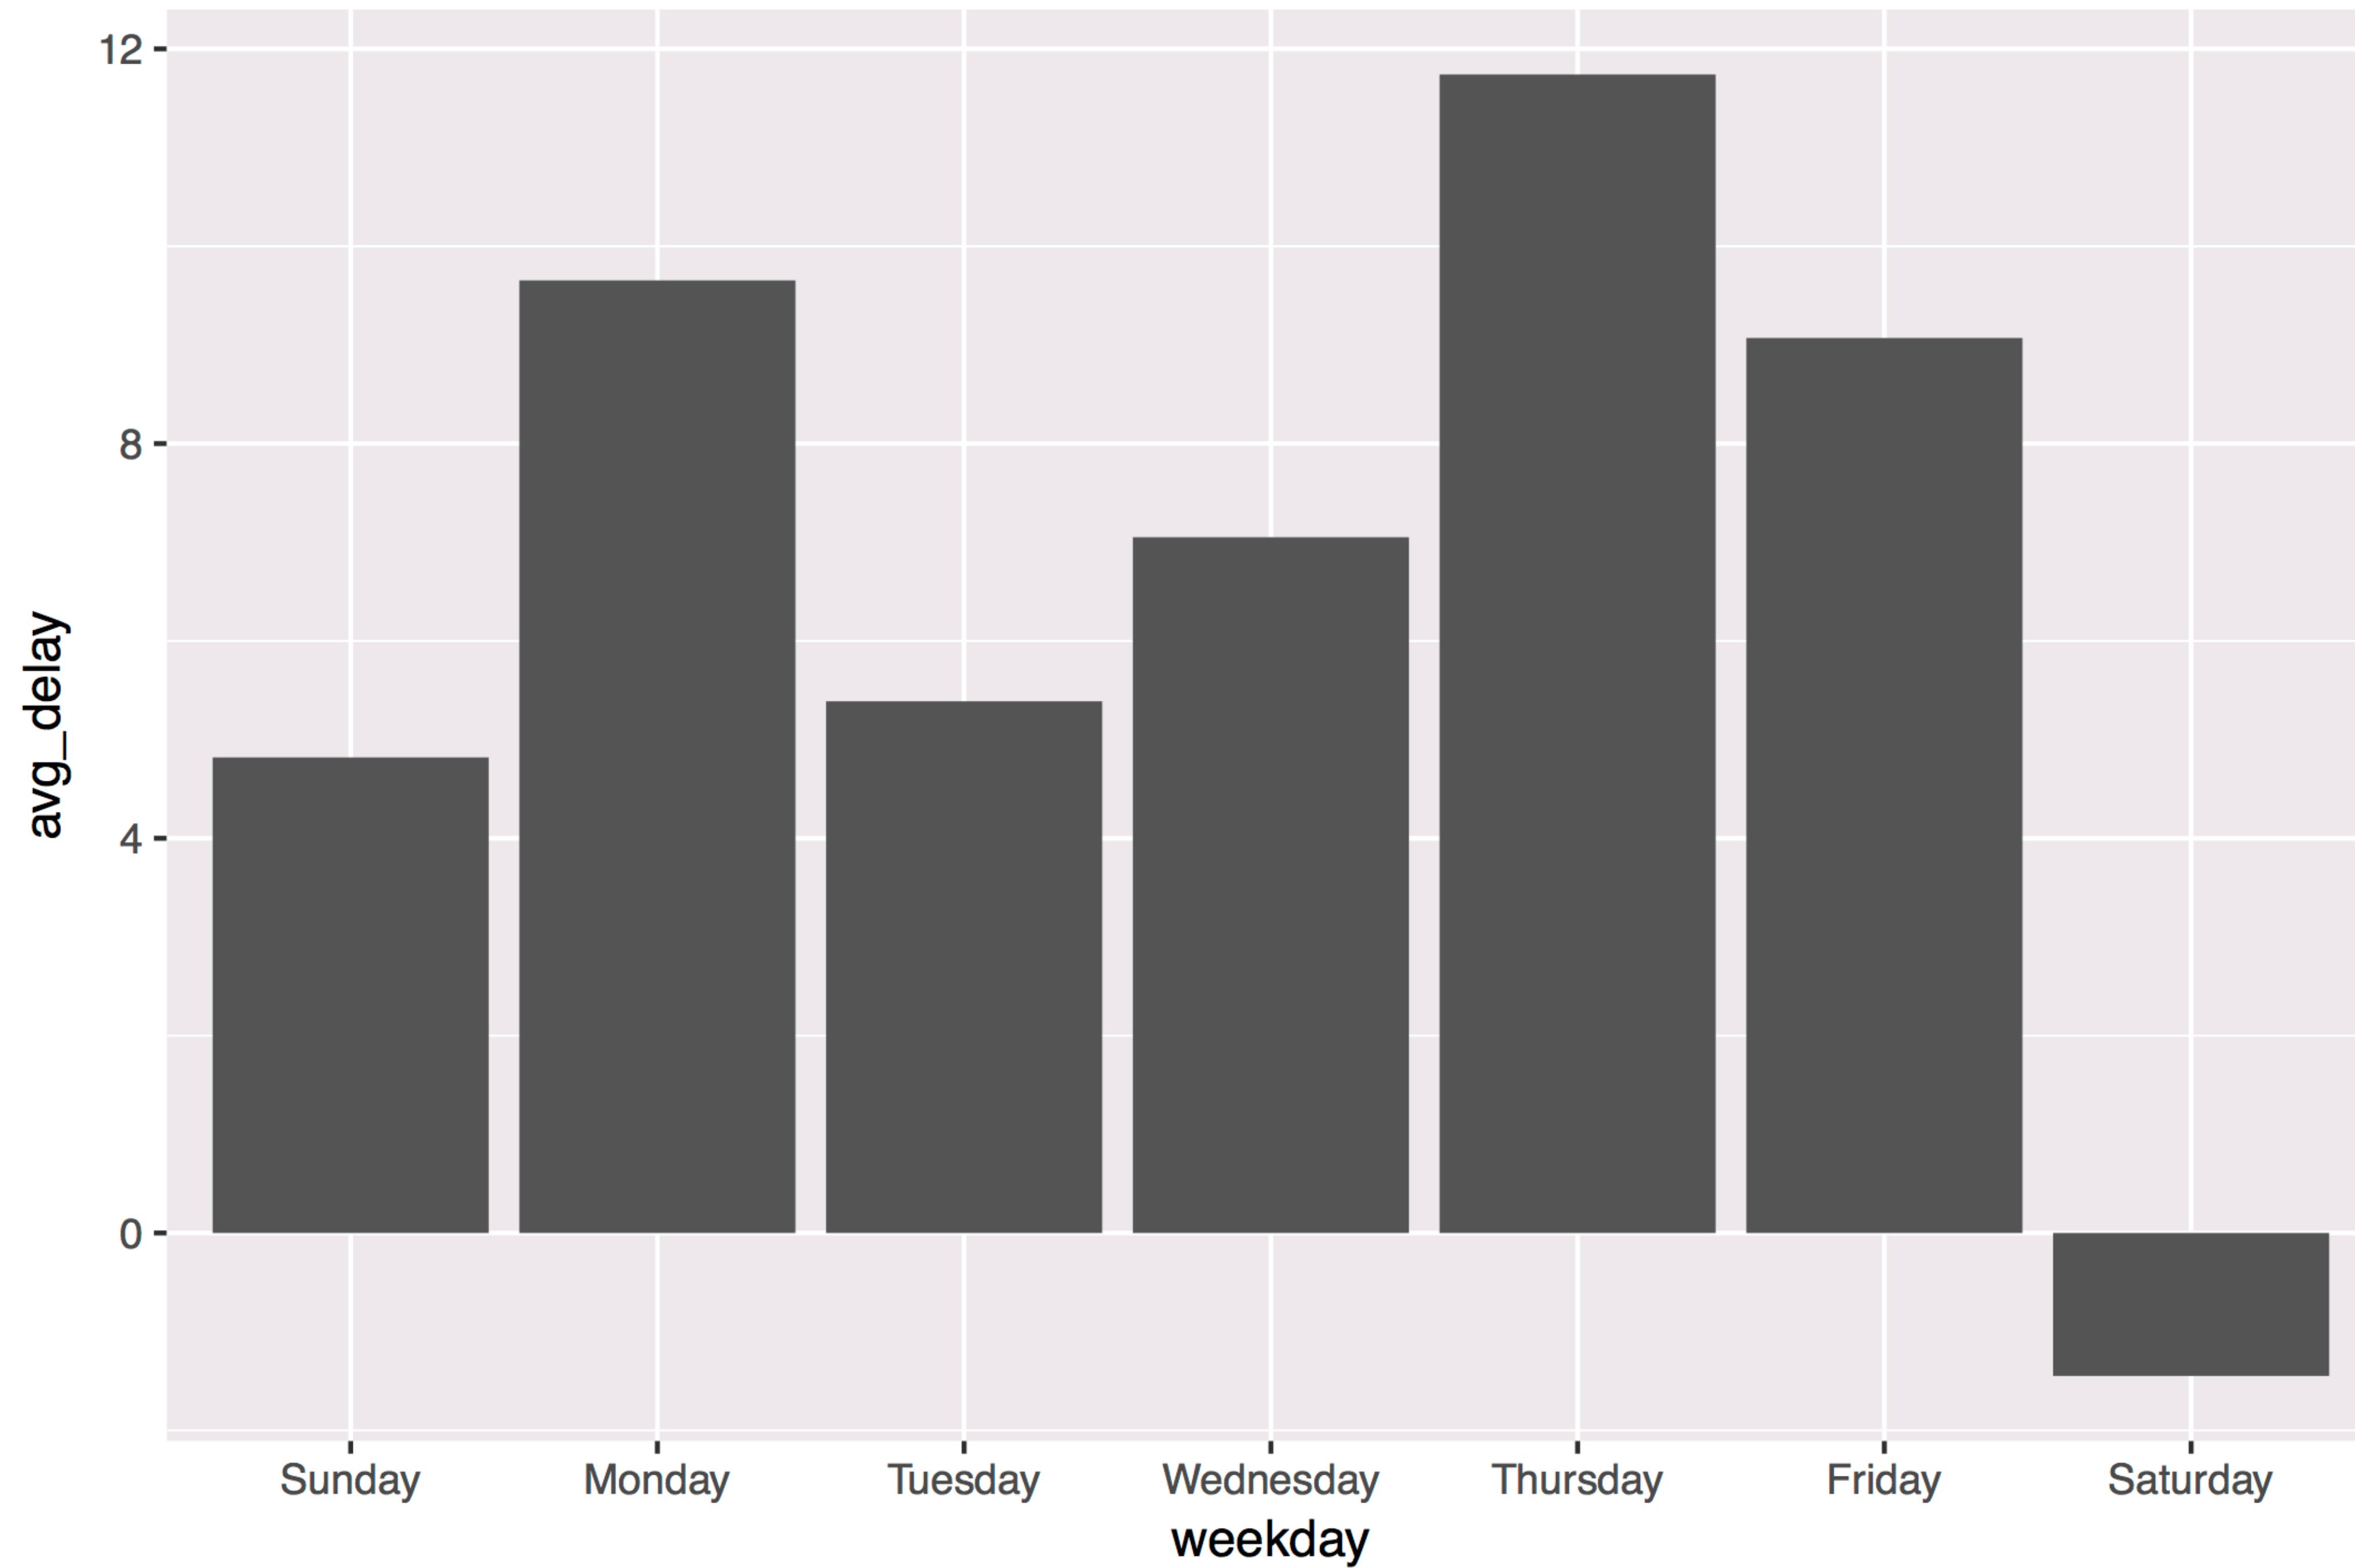
Calculate the average **arr_delay** by day of the week.

Plot the results as a column chart (bar chart) with **geom_col()**.



```
flights %>%  
  mutate(weekday = wday(time_hour, label = TRUE, abbr = FALSE)) %>%  
  group_by(weekday) %>%  
  drop_na(arr_delay) %>%  
  summarise(avg_delay = mean(arr_delay)) %>%  
  ggplot() +  
    geom_col(mapping = aes(x = weekday, y = avg_delay))
```





Data types with



Data Types

Main Ideas

Notes

Notes Form

Please Take out